

Customer Performance Analysis via Regression

You can find the full code, data, and documentation for this project on GitHub: CSCA-5622-Supervised-Learning-Final-Project (github.com/treinart) <https://github.com/treinart/CSCA-5622-Supervised-Learning-Final-Project.git>

This project uses supervised regression modeling to analyze dealership customer performance at the customer level. The goal is to predict each customer's average total sales, combining labor and parts, and to understand what characteristics make some customers more valuable than others.

Data Source and Description

The dataset was generated using a custom Python script (`generate_invoice_data.py`). It simulates dealership invoice activity using business-driven logic. All customer names, invoice numbers, and sales totals are synthetic to protect confidentiality.

The file `anonymized_invoice_data.csv` contains 45,514 invoices with 16 columns, including job type, location, labor and parts sales, and customer identifiers.

Project Dependencies and Environment

This notebook requires the following Python packages:

- pandas
- numpy
- matplotlib
- seaborn
- scikit-learn
- scipy
- xgboost

If you are running this notebook for the first time, please ensure all packages are installed in your Jupyter environment. You can install any missing packages using pip (run these commands in a Jupyter code cell or your terminal):

```
!pip install pandas numpy matplotlib seaborn scikit-learn scipy xgboost
```

If you see any import errors, check that your notebook kernel matches your Python environment.

To verify, run:

```
import sys print(sys.executable)
```

This tells you which Python Jupyter is using.

If all else fails From command prompt, run:

```
<full_path_from_above> -m pip install xgboost
```

This was the only way I was successful installing XGBoost. This will guarantee that XGBoost is installed in the environment Jupyter is using. Restart the Jupyter kernel after installing. In a code cell, try:

```
from xgboost import XGBRegressor
```

If you do not get an error, it worked.

```
In [1]: # Import core packages
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
# Set style for plots
sns.set_theme(style="darkgrid", palette="bright", context="talk")
```

Loading Invoice Data for Analysis

We now load the dealership invoice data into a pandas DataFrame for inspection and downstream processing.

```
In [2]: # Load the invoice-level dataset from CSV using pandas
df = pd.read_csv("anonymized_invoice_data.csv") # This loads the entire file into a DataFrame called df

# Show the first five rows to confirm successful loading and review column names and sample data
df.head()
```

```
Out[2]:
```

	Location	Whse	InvoiceNo	InvDate	CustNo	CustName	RORtype	Dept	HoursWorked	HoursBilled	LaborBilled\$	Labo
0	Dallas	DL2	100000	05/12/2025	JJQ9K	Goldstar Associates	COUNTER	30	0.00	0.00	0.00	
1	Green Bay	GB1	100001	08/25/2022	JJQ9K	Goldstar Associates	TRAILER	40	2.44	1.94	604.28	
2	Dallas	DL1	100002	09/21/2022	JJQ9K	Goldstar Associates	RESALE	40	6.27	7.01	1870.92	
3	Chicago	CH1	100003	04/29/2025	JJQ9K	Goldstar Associates	RESALE	20	5.40	7.40	1940.86	
4	Green Bay	GB4	100004	10/26/2023	JJQ9K	Goldstar Associates	TRAILER	50	4.68	5.23	1270.51	

Data Structure and Cleaning Checks

We now check the structure of the data, look for missing values, and confirm column types. These steps ensure that the data is ready for aggregation and modeling.

The generator script was designed to avoid missing or illogical values, but we confirm data integrity here. We check for missing values, review data types, and plot key fields for outliers or oddities.

```
In [3]: # Display the number of rows and columns in the dataset
df.shape # Output: (number of rows, number of columns)
print("Rows, columns:", df.shape)
```

Rows, columns: (45514, 15)

```
In [4]: # Print the number of unique customers in the dataset.
# This counts how many distinct customer IDs (CustNo) appear in the invoice-level data.
print("Number of unique customers:", df["CustNo"].nunique())
```

Number of unique customers: 387

```
In [5]: # Check for missing values in each column to confirm data integrity
df.isnull().sum() # Expect all columns to show 0 missing values
print("\nMissing values:\n", df.isnull().sum())
```

Missing values:
Location 0
Whse 0
InvoiceNo 0
InvDate 0
CustNo 0
CustName 0
ROtype 0
Dept 0
HoursWorked 0
HoursBilled 0
LaborBilled\$ 0
LaborWorked\$ 0
PartsSales\$ 0
PartsCost\$ 0
InvCycleDays 0
dtype: int64

```
In [6]: # Review the data types (e.g., float, int, object) for each column
df.dtypes # This helps identify if any columns are mis-typed
print("\nColumn types:\n", df.dtypes)
```

Column types:
Location object
Whse object
InvoiceNo int64
InvDate object
CustNo object
CustName object
ROtype object
Dept int64
HoursWorked float64
HoursBilled float64
LaborBilled\$ float64
LaborWorked\$ float64
PartsSales\$ float64
PartsCost\$ float64
InvCycleDays float64
dtype: object

```
In [7]: # Show summary statistics (mean, std, min, max, quartiles) for all numeric columns
df.describe()
```

Out[7]:

	InvoiceNo	Dept	HoursWorked	HoursBilled	LaborBilled\$	LaborWorked\$	PartsSales\$	PartsCost\$	Inv
count	45514.000000	45514.000000	45514.000000	45514.000000	45514.000000	45514.000000	45514.000000	45514.000000	45514.000000
mean	122756.500000	30.054489	1.943942	2.473220	902.099548	308.744873	828.727662	534.102356	122756.500000
std	13138.904413	14.095500	2.400496	3.195262	1342.479969	388.871755	520.255687	298.860757	13138.904413
min	100000.000000	10.000000	0.000000	0.000000	0.000000	0.000000	-299.360000	-127.570000	100000.000000
25%	111378.250000	20.000000	0.000000	0.000000	0.000000	0.000000	415.047500	281.980000	111378.250000
50%	122756.500000	30.000000	0.000000	0.000000	0.000000	0.000000	776.615000	527.845000	122756.500000
75%	134134.750000	40.000000	4.110000	4.940000	1614.620000	631.750000	1158.635000	776.700000	134134.750000
max	145513.000000	50.000000	6.900000	14.490000	13222.220000	1207.500000	2975.650000	1251.410000	145513.000000

Visualizing Labor Billed \$ Distribution

To better understand the distribution and possible outliers in the Labor Billed \$ field, we use three complementary visualizations:

- a zoomed-in box plot,
- a standard histogram,
- and a histogram with a logarithmic x-axis.

Main goal:

Quickly see where most jobs fall, whether there are lots of outliers, and whether the distribution is skewed (which can affect modeling and feature engineering).

Secondary goal:

Identify whether any data cleaning, scaling, or transformation might be required before modeling.

```
In [8]: # Filter for service jobs only
service_jobs = df[df['ROtype'].isin(['RESALE', 'TRUCK', 'TRAILER'])].copy()

# Compute quartiles for IQR
q1 = service_jobs["LaborBilled$"].quantile(0.25)
q3 = service_jobs["LaborBilled$"].quantile(0.75)
iqr = q3 - q1

median_val = service_jobs["LaborBilled$"].median()
mean_val = service_jobs["LaborBilled$"].mean()
outlier_count = (service_jobs["LaborBilled$"] > 5000).sum()
total_count = service_jobs["LaborBilled$"].count()

print(f"Labor Billed $ stats for RESALE, TRUCK, and TRAILER jobs only:")
print(f"  Median: ${median_val:,.0f}")
print(f"  Mean:   ${mean_val:,.0f}")
print(f"  Outliers above $5000: {outlier_count} out of {total_count}")

plt.figure(figsize=(12, 6), facecolor='black')
flierprops = dict(marker='o', markerfacecolor='crimson', markeredgecolor='white', markersize=8, linestyle='none', alpha=0.5)
ax = sns.boxplot(
    x=service_jobs["LaborBilled$"],
    color='dodgerblue',
    fliersize=6, linewidth=3,
    flierprops=flierprops
)
plt.title("Labor Billed $ Boxplot (Service Jobs Only)", fontsize=26, color='darkorange', fontweight='bold', pad=10)
plt.xlabel("Labor Billed $", fontsize=22, color='darkorange', fontweight='bold', labelpad=10)
plt.xlim(0, 8000)
plt.xticks(fontsize=14, color='white')
plt.yticks([])
ax.set_facecolor('black')

# Add horizontal grid for readability
ax.xaxis.grid(True, color='gray', linestyle='-', linewidth=0.75, alpha=0.45)

for spine in ax.spines.values():
    spine.set_edgecolor('darkorange')
    spine.set_linewidth(2.7)

bbox_props = dict(boxstyle="round,pad=0.5", fc="black", ec="yellow", lw=2.7, alpha=0.80)
plt.text(0.98, 0.05,
        f"Median: ${median_val:,.0f}\nMean: ${mean_val:,.0f}\nOutliers above $5000: {outlier_count}",
        ha='right', va='bottom', fontsize=17, color='yellow', fontweight='bold',
        transform=ax.transAxes, bbox=bbox_props)

plt.tight_layout()
plt.show()
```

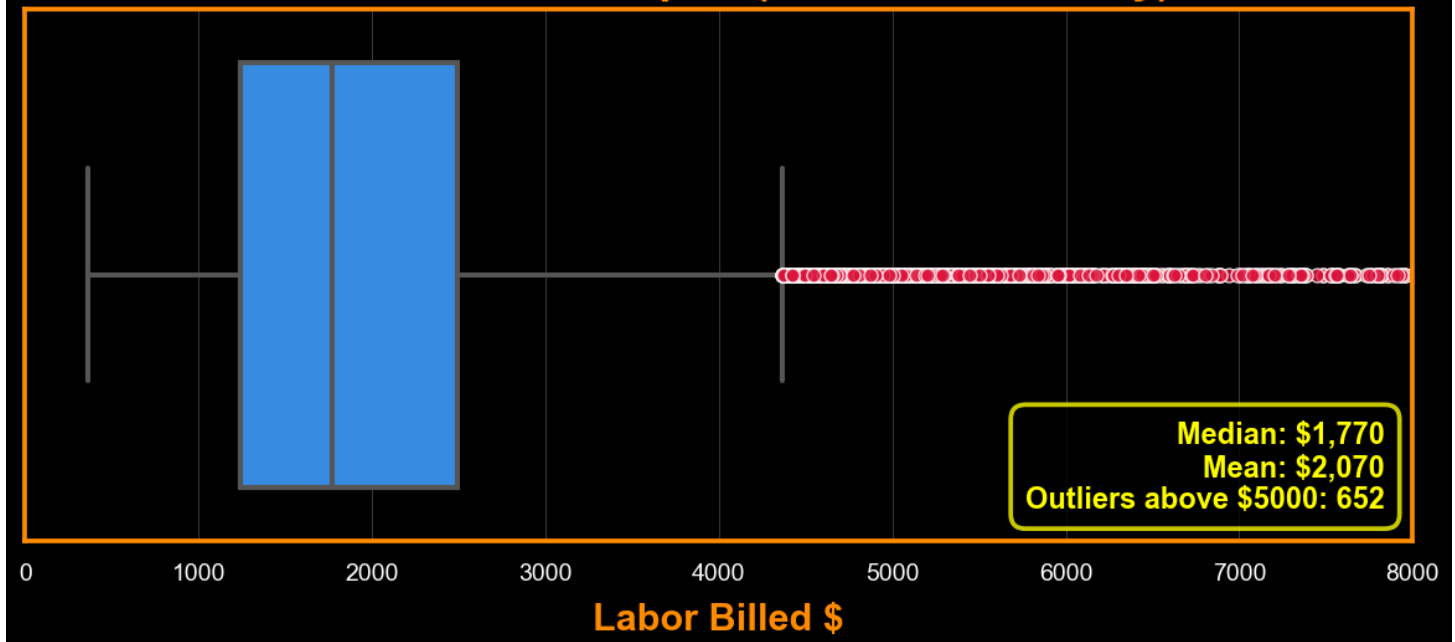
Labor Billed \$ stats for RESALE, TRUCK, and TRAILER jobs only:

Median: \$1,770

Mean: \$2,070

Outliers above \$5000: 652 out of 19839

Labor Billed \$ Boxplot (Service Jobs Only)



```
In [9]: plt.figure(figsize=(13, 7), facecolor='black')
flierprops = dict(marker='o', markerfacecolor='lime', markeredgecolor='crimson', markersize=10, linestyle='none', alpha=0.5)
ax = sns.boxplot(
    x="R0type", y="LaborBilled$",
    data=service_jobs,
    order=['RESALE', 'TRAILER', 'TRUCK'],
    color='dodgerblue',
    fliersize=6, linewidth=2.2,
    flierprops=flierprops
)
plt.title("Labor Billed $ by R0type", fontsize=26, color='darkorange', fontweight='bold', pad=10)
plt.xlabel("R0type", fontsize=22, color='darkorange', fontweight='bold')
plt.ylabel("Labor Billed $", fontsize=22, color='darkorange', fontweight='bold')
plt.xticks(fontsize=20, color='white', fontweight='bold')
plt.yticks(fontsize=14, color='white')
ax.set_facecolor('black')
ax.xaxis.grid(True, color='gray', linestyle='--', linewidth=0.75, alpha=0.45)
for spine in ax.spines.values():
    spine.set_edgecolor('darkorange')
    spine.set_linewidth(2.2)


# Print summary stats for LaborBilled$ by R0type
group_stats = service_jobs.groupby("R0type")["LaborBilled$"].agg(["count", "median", "mean", "min", "max"])
group_stats = group_stats.loc[["RESALE", "TRAILER", "TRUCK"]] # Enforce order

print("Labor Billed $ summary by R0type:")
for ro in group_stats.index:
    row = group_stats.loc[ro]
    print(f" {ro:8s} | Count: {int(row['count']):5d}  Median: ${row['median']:7,.0f}  Mean: ${row['mean']:7,.0f}  Min: ${row['min']:7,.0f}  Max: ${row['max']:7,.0f}")

plt.tight_layout()
plt.show()
```

Labor Billed \$ summary by R0type:

RESALE	Count: 13041	Median: \$ 1,974	Mean: \$ 2,272	Min: \$533	Max: \$13,222
TRAILER	Count: 5815	Median: \$ 1,362	Mean: \$ 1,572	Min: \$358	Max: \$9,187
TRUCK	Count: 983	Median: \$ 2,041	Mean: \$ 2,336	Min: \$585	Max: \$12,250



```
In [10]: # Calculate medians to order locations by typical labor billed
medians = service_jobs.groupby("Location")["LaborBilled$"].median().sort_values(ascending=False)
ordered_locations = medians.index.tolist()

plt.figure(figsize=(16, 8), facecolor='black')
flierprops = dict(marker='o', markerfacecolor='lime', markeredgecolor='crimson', markersize=8, linestyle='none', alpha=0.5)
ax = sns.boxplot(
    x="Location", y="LaborBilled$",
    data=service_jobs,
    order=ordered_locations,
    color='dodgerblue',
    fliersize=5, linewidth=2.1,
    flierprops=flierprops
)


plt.title("Labor Billed $ by Location", fontsize=26, color='darkorange', fontweight='bold', pad=10)
plt.xlabel("Location", fontsize=22, color='darkorange', fontweight='bold')
plt.ylabel("Labor Billed $", fontsize=22, color='darkorange', fontweight='bold')
plt.xticks(fontsize=15, color='white', fontweight='bold', rotation=30)
plt.yticks(fontsize=14, color='white')
ax.set_facecolor('black')
ax.xaxis.grid(True, color='gray', linestyle='--', linewidth=0.75, alpha=0.45)
for spine in ax.spines.values():
    spine.set_edgecolor('darkorange')
    spine.set_linewidth(2.2)

# Print summary stats for LaborBilled$ by Location (top 5 by median)
location_stats = service_jobs.groupby("Location")["LaborBilled$"].agg(["count", "median", "mean", "min", "max"])
top_locations = medians.index[:5] # Show top 5 Locations by median Labor billed

print("Labor Billed $ summary for top 5 Locations (by median):")
for loc in top_locations:
    row = location_stats.loc[loc]
    print(f" {loc:10s} | Count: {int(row['count']):5d}  Median: ${row['median']:7,.0f}  Mean: ${row['mean']:7,.0f}

plt.tight_layout()
plt.show()
```

```
Labor Billed $ summary for top 5 Locations (by median):
Chicago      | Count: 3203  Median: $ 2,196  Mean: $ 2,612  Min: $439  Max: $13,222
Dallas       | Count: 3288  Median: $ 2,006  Mean: $ 2,470  Min: $365  Max: $13,183
Green Bay    | Count: 3264  Median: $ 1,876  Mean: $ 2,151  Min: $403  Max: $12,250
New York     | Count: 3367  Median: $ 1,630  Mean: $ 1,780  Min: $404  Max: $8,690
Los Angeles  | Count: 3341  Median: $ 1,586  Mean: $ 1,743  Min: $358  Max: $8,258
```



```
In [11]: # Compute total sales per invoice
df['TotalSales$'] = df['LaborBilled$'] + df['PartsSales$']

# Aggregate total sales by customer
customer_sales = df.groupby('CustName')['TotalSales$'].sum().sort_values(ascending=False)
top_customers = customer_sales.head(10).index.tolist()

# Filter for invoices from top 10 customers
top_cust_df = df[df['CustName'].isin(top_customers)].copy()


# Enforce order (so boxplot is sorted by total sales, not alphabetically)
top_cust_df['CustName'] = pd.Categorical(top_cust_df['CustName'], categories=top_customers, ordered=True)

# Print summary stats for the top 10 customers
cust_stats = top_cust_df.groupby("CustName", observed=True)["TotalSales$"].agg(["count", "median", "mean", "min", "max"])
print("Total Sales summary for Top 10 Customers:")
for cust in top_customers:
    row = cust_stats.loc[cust]
    print(f" {cust:25s} | Count: {int(row['count']):5d} Median: ${row['median']:7,.0f} Mean: ${row['mean']:7,.0f}")

# Create the boxplot
plt.figure(figsize=(15, 8), facecolor='black')
flierprops = dict(marker='o', markerfacecolor='lime', markeredgecolor='crimson', markersize=8, linestyle='none', alpha=0.5)
ax = sns.boxplot(
    x="CustName", y="TotalSales$",
    data=top_cust_df,
    order=top_customers,
    color='dodgerblue',
    fliersize=6, linewidth=2.2,
    flierprops=flierprops
)
plt.title("Total Sales per Invoice for Top 10 Customers", fontsize=24, color='darkorange', fontweight='bold', pad=10)
plt.xlabel("Customer", fontsize=19, color='darkorange', fontweight='bold')
plt.ylabel("Total Sales per Invoice ($)", fontsize=19, color='darkorange', fontweight='bold')
plt.xticks(fontsize=13, color='white', fontweight='bold', rotation=28, ha='right')
plt.yticks(fontsize=14, color='white')
ax.set_facecolor('black')
ax.xaxis.grid(True, color='gray', linestyle='-', linewidth=0.75, alpha=0.45)
for spine in ax.spines.values():
    spine.set_edgecolor('darkorange')
    spine.set_linewidth(2.2)
plt.tight_layout()
plt.show()
```

Total Sales summary for Top 10 Customers:

Paramount Transfer	Count: 2008	Median: \$ 1,738	Mean: \$ 1,927	Min: \$38	Max: \$11,089
Paramount Logistics	Count: 1689	Median: \$ 1,856	Mean: \$ 2,007	Min: \$58	Max: \$9,951
Acme Inc	Count: 1747	Median: \$ 1,587	Mean: \$ 1,814	Min: \$76	Max: \$7,922
Liberty Industries	Count: 1677	Median: \$ 1,557	Mean: \$ 1,867	Min: \$48	Max: \$11,336
Stonegate Partners	Count: 1386	Median: \$ 1,788	Mean: \$ 2,113	Min: \$36	Max: \$13,233
Goldstar Transfer	Count: 1522	Median: \$ 1,534	Mean: \$ 1,653	Min: \$23	Max: \$6,708
Unity Resources	Count: 1295	Median: \$ 1,366	Mean: \$ 1,797	Min: \$19	Max: \$10,359
Aurora Inc	Count: 1245	Median: \$ 1,511	Mean: \$ 1,580	Min: \$-35	Max: \$5,605
Acme Systems	Count: 842	Median: \$ 1,989	Mean: \$ 2,194	Min: \$80	Max: \$9,157
Aurora Group	Count: 853	Median: \$ 1,588	Mean: \$ 2,099	Min: \$76	Max: \$12,093



```
In [12]: # Filter to only service jobs for Labor sales analysis
service_jobs = df[df['ROtype'].isin(['RESALE', 'TRUCK', 'TRAILER'])].copy()

# Print summary stats for Labor Billed $ (service jobs only)
median_val = service_jobs["LaborBilled$"].median()
mean_val = service_jobs["LaborBilled$"].mean()
min_val = service_jobs["LaborBilled$"].min()
max_val = service_jobs["LaborBilled$"].max()
above_5000 = (service_jobs["LaborBilled$"] > 5000).sum()
total_count = service_jobs["LaborBilled$"].count()

print("Labor Billed $ summary (service jobs only):")
print(f" Median: ${median_val:,.0f}")
print(f" Mean:  ${mean_val:,.0f}")
print(f" Min:    ${min_val:,.0f}")
print(f" Max:    ${max_val:,.0f}")
print(f" Invoices above $5,000: {above_5000} out of {total_count}")

# Improved Histogram for Labor Billed $ (Service Jobs)
plt.figure(figsize=(14, 7), facecolor='black')
ax = sns.histplot(
    service_jobs["LaborBilled$"],
    bins=120,
    kde=True,
    color='dodgerblue',
    edgecolor='white',
    alpha=0.86
)

plt.title("Distribution of Labor Billed $ (Service Jobs Only)", fontsize=25, color='darkorange', fontweight='bold',
plt.xlabel("Labor Billed $", fontsize=22, color='darkorange', fontweight='bold')
plt.ylabel("Number of Invoices", fontsize=22, color='darkorange', fontweight='bold')
plt.xticks(fontsize=16, color='white')
plt.yticks(fontsize=16, color='white')
```



```

ax.set_facecolor('black')
for spine in ax.spines.values():
    spine.set_edgecolor('darkorange')
    spine.set_linewidth(2.1)
ax.xaxis.grid(True, color='gray', linestyle=':', linewidth=0.7, alpha=0.43)


# Median and mean Lines
plt.axvline(median_val, color='yellow', linestyle='--', linewidth=3, label=f"Median (${median_val:,.0f})")
plt.axvline(mean_val, color='crimson', linestyle='-', linewidth=3, label=f"Mean (${mean_val:,.0f})")
leg = plt.legend(fontsize=19, loc='upper right', facecolor='black', edgecolor='yellow', frameon=True)
for text in leg.get_texts():
    text.set_color("white")

plt.tight_layout()
plt.show()

```

Labor Billed \$ summary (service jobs only):

Median:	\$1,770
Mean:	\$2,070
Min:	\$358
Max:	\$13,222
Invoices above \$5,000:	652 out of 19839



```

In [13]: # Filter to only service jobs for Labor sales analysis
service_jobs = df[df['R0type'].isin(['RESALE', 'TRUCK', 'TRAILER'])].copy()

# Print summary stats for Labor Billed $ (service jobs only)
median_val = service_jobs["LaborBilled$"].median()
mean_val = service_jobs["LaborBilled$"].mean()
min_val = service_jobs["LaborBilled$"].min()
max_val = service_jobs["LaborBilled$"].max()
above_5000 = (service_jobs["LaborBilled$"] > 5000).sum()
total_count = service_jobs["LaborBilled$"].count()

print("Labor Billed $ summary (service jobs only):")
print(f"  Median: ${median_val:,.0f}")
print(f"  Mean:   ${mean_val:,.0f}")
print(f"  Min:    ${min_val:,.0f}")
print(f"  Max:    ${max_val:,.0f}")
print(f"  Invoices above $5,000: {above_5000} out of {total_count}")

# Histogram with Logarithmic X-Axis
plt.figure(figsize=(14, 7), facecolor='black')
ax = sns.histplot(
    service_jobs["LaborBilled$"],
    bins=120,
    kde=True,
    color='dodgerblue',

```

```

edgecolor='white',
alpha=0.86,
log_scale=(True, False) # Logarithmic x-axis, normal y-axis
)

plt.title("Distribution of Labor Billed $ (Service Jobs Only, Log Scale)", fontsize=25, color='darkorange', fontweight='bold')
plt.xlabel("Labor Billed $ (Log Scale)", fontsize=22, color='darkorange', fontweight='bold')
plt.ylabel("Number of Invoices", fontsize=22, color='darkorange', fontweight='bold')
plt.xticks(fontsize=16, color='white')
plt.yticks(fontsize=16, color='white')
ax.set_facecolor('black')
for spine in ax.spines.values():
    spine.set_edgecolor('darkorange')
    spine.set_linewidth(2.1)
ax.xaxis.grid(True, color='gray', linestyle=':', linewidth=0.7, alpha=0.43)

# Median and mean Lines
plt.axvline(median_val, color='yellow', linestyle='--', linewidth=3, label=f"Median (${median_val:,.0f})")
plt.axvline(mean_val, color='crimson', linestyle='-', linewidth=3, label=f"Mean (${mean_val:,.0f})")
leg = plt.legend(fontsize=19, loc='upper right', facecolor='black', edgecolor='yellow', frameon=True)
for text in leg.get_texts():
    text.set_color("white")

plt.tight_layout()
plt.show()


```

Labor Billed \$ summary (service jobs only):

```

Median: $1,770
Mean:   $2,070
Min:    $358
Max:    $13,222
Invoices above $5,000: 652 out of 19839

```



```

In [14]: # Filter for only service jobs if you want, or use all jobs:
# parts_jobs = df[df['ROtype'].isin(['RESALE', 'TRUCK', 'TRAILER'])].copy()
parts_jobs = df.copy()

# Print summary stats for Parts Sales $ (all invoices)
median_val = parts_jobs["PartsSales$"].median()
mean_val = parts_jobs["PartsSales$"].mean()
min_val = parts_jobs["PartsSales$"].min()
max_val = parts_jobs["PartsSales$"].max()
above_2000 = (parts_jobs["PartsSales$"] > 2000).sum()
total_count = parts_jobs["PartsSales$"].count()

print("Parts Sales $ summary (all invoices):")
print(f"  Median: ${median_val:,.0f}")
print(f"  Mean:   ${mean_val:,.0f}")

```

```

print(f"   Min:    ${min_val:,.0f}")
print(f"   Max:    ${max_val:,.0f}")
print(f"   Invoices above $2,000: {above_2000} out of {total_count}")

# Improved Histogram for Parts Sales $
plt.figure(figsize=(14, 7), facecolor='black')
ax = sns.histplot(
    parts_jobs["PartsSales$"],
    bins=120,
    kde=True,
    color='limegreen',
    edgecolor='crimson',
    alpha=0.84
)

plt.title("Distribution of Parts Sales $ (All Invoices)", fontsize=25, color='darkorange', fontweight='bold', pad=10)
plt.xlabel("Parts Sales $", fontsize=22, color='darkorange', fontweight='bold')
plt.ylabel("Number of Invoices", fontsize=22, color='darkorange', fontweight='bold')
plt.xticks(fontsize=16, color='white')
plt.yticks(fontsize=16, color='white')
ax.set_facecolor('black')
for spine in ax.spines.values():
    spine.set_edgecolor('darkorange')
    spine.set_linewidth(2.1)
ax.xaxis.grid(True, color='gray', linestyle=':', linewidth=0.7, alpha=0.43)

# Median and mean Lines
plt.axvline(median_val, color='yellow', linestyle='--', linewidth=3, label=f"Median (${median_val:,.0f})")
plt.axvline(mean_val, color='magenta', linestyle='-', linewidth=3, label=f"Mean (${mean_val:,.0f})")
leg = plt.legend(fontsize=19, loc='upper right', facecolor='black', edgecolor='yellow', frameon=True)
for text in leg.get_texts():
    text.set_color("white")

plt.tight_layout()
plt.show()


```

Parts Sales \$ summary (all invoices):

```

Median: $777
Mean:   $829
Min:    $-299
Max:    $2,976
Invoices above $2,000: 1183 out of 45514

```



Aggregate to Customer Level

We create a customer-level table by summarizing each customer's invoices.

```
In [15]: # Calculate total sales for each invoice
df["TotalSales$"] = df["LaborBilled$"] + df["PartsSales$"]

# Calculate labor and parts gross margin percent (as a decimal)
df["LaborGM%"] = 1 - (df["LaborWorked$"] / df["LaborBilled$"])
df["PartsGM%"] = 1 - (df["PartsCost$"] / df["PartsSales$"])

# Calculate efficiency (hours billed divided by hours worked)
df["Efficiency"] = df["HoursBilled"] / df["HoursWorked"]

# Group by customer to create customer-level features
customer_df = df.groupby(["CustNo", "CustName"]).agg(
    AvgTotalSalesPerInvoice=("TotalSales$", "mean"),
    TotalInvoices=("InvoiceNo", "count"),
    AvgLaborGM=("LaborGM%", "mean"),
    AvgPartsGM=("PartsGM%", "mean"),
    AvgEfficiency=("Efficiency", "mean"),
    MostCommonROtype=("ROtype", lambda x: x.mode()[0] if not x.mode().empty else 'UNKNOWN'),
    MostCommonDept=("Dept", lambda x: x.mode()[0] if not x.mode().empty else -1),
    MostCommonLocation=("Location", lambda x: x.mode()[0] if not x.mode().empty else 'UNKNOWN')
).reset_index()

customer_df.head()
```

Out[15]:

	CustNo	CustName	AvgTotalSalesPerInvoice	TotalInvoices	AvgLaborGM	AvgPartsGM	AvgEfficiency	MostCommonROtype
0	016ZZ	Goldstar Solutions	2005.883652	742	0.658874	0.345703	1.370037	COUNTER
1	04CDPQ	Summit Partners	1705.610000	23	0.574704	0.223554	1.024120	COUNTER
2	050E6	Evergreen Systems	2176.141463	82	0.611540	0.232640	1.076672	RESALE
3	0AOBT3	Unity Logistics	1369.023846	13	0.592245	0.323652	1.501356	COUNTER
4	0KB68	Liberty Solutions	1885.920909	11	0.756308	0.335894	1.298439	COUNTER

Customer-Level EDA

We examine the aggregated customer dataset to understand feature distributions, relationships, and trends that might inform modeling.

```
In [16]: # Print the number of unique customers in the dataset.
print("Number of unique customers:", df["CustNo"].nunique())
# Show basic stats and check shape (should match your customer count)
print("Customer, Columns", customer_df.shape)
customer_df.describe()
```

Number of unique customers: 387
Customer, Columns (387, 10)

Out[16]:

	AvgTotalSalesPerInvoice	TotalInvoices	AvgLaborGM	AvgPartsGM	AvgEfficiency	MostCommonDept
count	387.000000	387.000000	375.000000	387.000000	375.000000	387.000000
mean	1729.194152	117.607235	0.611134	0.311061	1.236275	27.235142
std	573.425260	276.293905	0.086646	0.065597	0.218668	13.422354
min	86.000000	1.000000	0.330004	0.109550	0.597656	10.000000
25%	1385.931938	13.000000	0.556711	0.275666	1.107004	20.000000
50%	1641.388378	28.000000	0.617406	0.305224	1.236414	30.000000
75%	1996.960852	68.500000	0.662813	0.350300	1.380155	40.000000
max	5563.240000	2008.000000	0.837117	0.610001	1.841010	50.000000

In [17]:

```
plt.figure(figsize=(13, 7), facecolor='black')

# Print summary statistics for AvgTotalSalesPerInvoice at the customer level
median_val = customer_df["AvgTotalSalesPerInvoice"].median()
mean_val = customer_df["AvgTotalSalesPerInvoice"].mean()
min_val = customer_df["AvgTotalSalesPerInvoice"].min()
max_val = customer_df["AvgTotalSalesPerInvoice"].max()
q25 = customer_df["AvgTotalSalesPerInvoice"].quantile(0.25)
q75 = customer_df["AvgTotalSalesPerInvoice"].quantile(0.75)

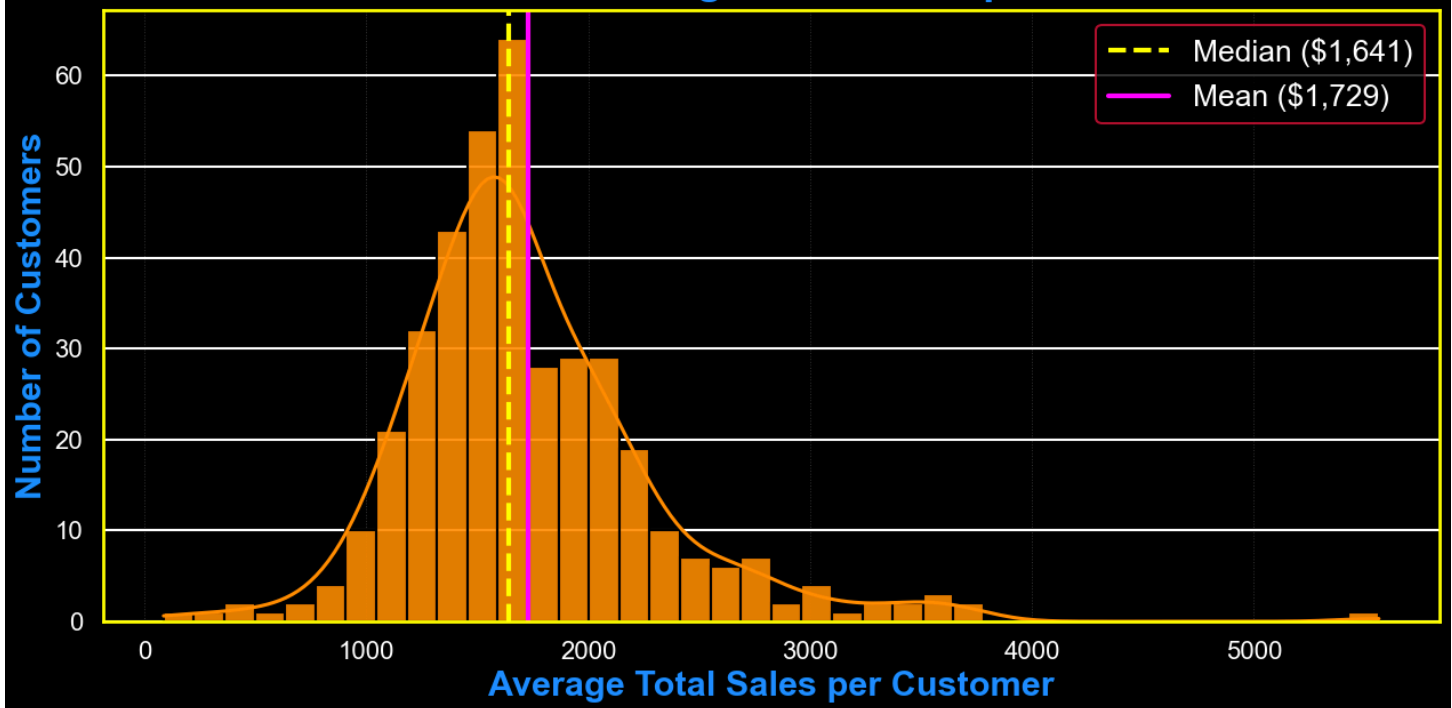
print("Average Total Sales per Customer summary:")
print(f"  Median: ${median_val:,.0f}")
print(f"  Mean:   ${mean_val:,.0f}")
print(f"  Min:    ${min_val:,.0f}")
print(f"  Max:    ${max_val:,.0f}")
print(f"  25th percentile: ${q25:,.0f}")
print(f"  75th percentile: ${q75:,.0f}")

ax = sns.histplot(customer_df["AvgTotalSalesPerInvoice"], bins=40, kde=True, color='darkorange', edgecolor='black', ax=
median_val = customer_df["AvgTotalSalesPerInvoice"].median()
mean_val = customer_df["AvgTotalSalesPerInvoice"].mean()
plt.axvline(median_val, color='yellow', linestyle='--', linewidth=3, label=f"Median (${median_val:,.0f})")
plt.axvline(mean_val, color='magenta', linestyle='-', linewidth=3, label=f"Mean (${mean_val:,.0f})")
plt.title("Distribution of Average Total Sales per Customer", fontsize=26, color='dodgerblue', fontweight='bold', pad
plt.xlabel("Average Total Sales per Customer", fontsize=22, color='dodgerblue', fontweight='bold')
plt.ylabel("Number of Customers", fontsize=22, color='dodgerblue', fontweight='bold')
plt.xticks(fontsize=16, color='white')
plt.yticks(fontsize=16, color='white')
ax.set_facecolor('black')
for spine in ax.spines.values():
    spine.set_edgecolor('yellow')
    spine.set_linewidth(2)
ax.xaxis.grid(True, color='gray', linestyle=':', linewidth=0.7, alpha=0.43)
leg = plt.legend(fontsize=19, loc='upper right', facecolor='black', edgecolor='crimson', frameon=True)
for text in leg.get_texts():
    text.set_color("white")
plt.tight_layout()
plt.show()
```

Average Total Sales per Customer summary:

Median: \$1,641
Mean: \$1,729
Min: \$86
Max: \$5,563
25th percentile: \$1,386
75th percentile: \$1,997

Distribution of Average Total Sales per Customer



```
In [18]: plt.figure(figsize=(13, 5), facecolor='black')

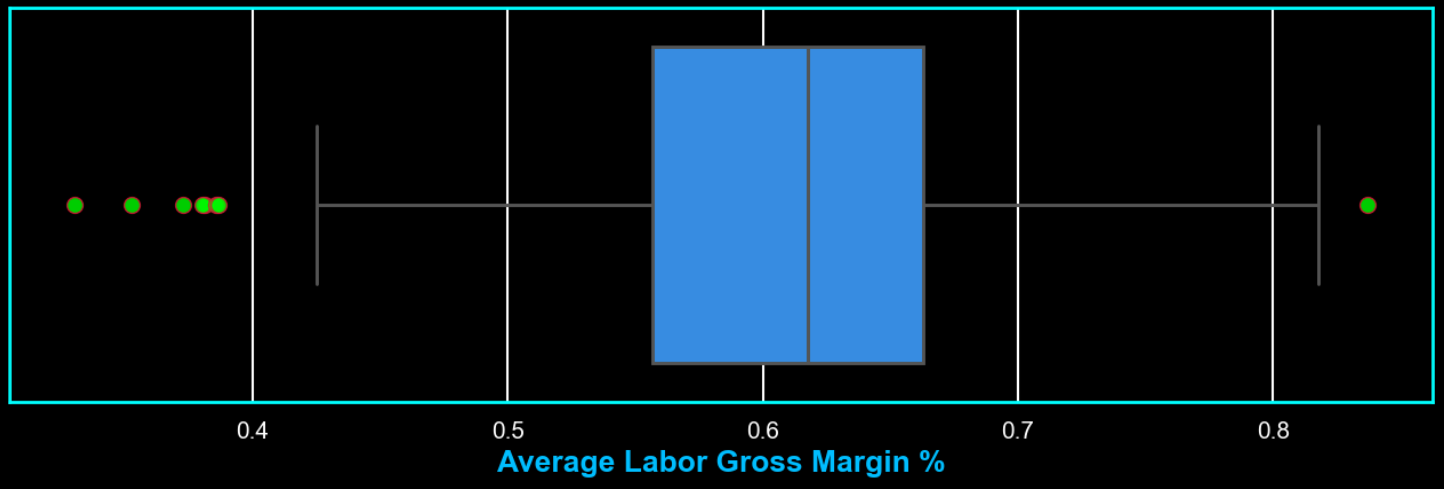
# Print summary statistics for average Labor gross margin percent per customer
median_gm = customer_df["AvgLaborGM"].median()
mean_gm = customer_df["AvgLaborGM"].mean()
min_gm = customer_df["AvgLaborGM"].min()
max_gm = customer_df["AvgLaborGM"].max()
p25_gm = customer_df["AvgLaborGM"].quantile(0.25)
p75_gm = customer_df["AvgLaborGM"].quantile(0.75)
print("Average Labor Gross Margin % per Customer summary:")
print(f"  Median: {median_gm:.2%}")
print(f"  Mean:   {mean_gm:.2%}")
print(f"  Min:    {min_gm:.2%}")
print(f"  Max:    {max_gm:.2%}")
print(f"  25th percentile: {p25_gm:.2%}")
print(f"  75th percentile: {p75_gm:.2%}")

ax = sns.boxplot(x=customer_df["AvgLaborGM"], color='dodgerblue', fliersize=7, linewidth=2.1,
                 flierprops=dict(marker='o', markerfacecolor='lime', markeredgecolor='crimson', markersize=10, linestyle='solid'))
plt.title("Boxplot of Average Labor Gross Margin % per Customer", fontsize=22, color='cyan', fontweight='bold', pad=10)
plt.xlabel("Average Labor Gross Margin %", fontsize=19, color='deepskyblue', fontweight='bold')
plt.xticks(fontsize=15, color='white')
plt.yticks([])
ax.set_facecolor('black')
for spine in ax.spines.values():
    spine.set_edgecolor('cyan')
    spine.set_linewidth(2)
plt.tight_layout()
plt.show()
```

Average Labor Gross Margin % per Customer summary:

```
Median: 61.74%
Mean:   61.11%
Min:    33.00%
Max:    83.71%
25th percentile: 55.67%
75th percentile: 66.28%
```

Boxplot of Average Labor Gross Margin % per Customer



```
In [19]: plt.figure(figsize=(13, 7), facecolor='black')

# Print summary statistics for AvgEfficiency per customer
eff_median = customer_df["AvgEfficiency"].median()
eff_mean = customer_df["AvgEfficiency"].mean()
eff_min = customer_df["AvgEfficiency"].min()
eff_max = customer_df["AvgEfficiency"].max()
eff_25 = customer_df["AvgEfficiency"].quantile(0.25)
eff_75 = customer_df["AvgEfficiency"].quantile(0.75)


print("Average Efficiency per Customer summary:")
print(f"  Median: {eff_median:.2f}")
print(f"  Mean:   {eff_mean:.2f}")
print(f"  Min:    {eff_min:.2f}")
print(f"  Max:     {eff_max:.2f}")
print(f"  25th percentile: {eff_25:.2f}")
print(f"  75th percentile: {eff_75:.2f}")

ax = sns.histplot(customer_df["AvgEfficiency"], bins=30, kde=True, color='darkorange', edgecolor='black', alpha=0.8)
median_val = customer_df["AvgEfficiency"].median()
mean_val = customer_df["AvgEfficiency"].mean()
plt.axvline(median_val, color='yellow', linestyle='--', linewidth=3, label=f"Median ({median_val:.2f})")
plt.axvline(mean_val, color='magenta', linestyle='-', linewidth=3, label=f"Mean ({mean_val:.2f})")
plt.title("Distribution of Average Efficiency per Customer", fontsize=26, color='crimson', fontweight='bold', pad=10)
plt.xlabel("Average Efficiency", fontsize=22, color='crimson', fontweight='bold')
plt.ylabel("Number of Customers", fontsize=22, color='crimson', fontweight='bold')
plt.xticks(fontsize=16, color='white')
plt.yticks(fontsize=16, color='white')
ax.set_facecolor('black')
for spine in ax.spines.values():
    spine.set_edgecolor('dodgerblue')
    spine.set_linewidth(2)
ax.xaxis.grid(True, color='gray', linestyle=':', linewidth=0.7, alpha=0.43)
leg = plt.legend(fontsize=19, loc='upper right', facecolor='black', edgecolor='lime', frameon=True)
for text in leg.get_texts():
    text.set_color("white")
plt.tight_layout()
plt.show()
```

Average Efficiency per Customer summary:

Median:	1.24
Mean:	1.24
Min:	0.60
Max:	1.84
25th percentile:	1.11
75th percentile:	1.38

Distribution of Average Efficiency per Customer



```
In [20]: plt.figure(figsize=(11, 9), facecolor='black')

numeric_cols = [
    "AvgTotalSalesPerInvoice",
    "TotalInvoices",
    "AvgLaborGM",
    "AvgPartsGM",
    "AvgEfficiency"
]
corr_matrix = customer_df[numeric_cols].corr()

print("Correlation Matrix Summary:")
for row in corr_matrix.index:
    for col in corr_matrix.columns:
        if row != col:
            print(f" {row:22s} vs. {col:22s}: {corr_matrix.loc[row, col]:.2f}")
print("-" * 50)

corr = customer_df[["AvgTotalSalesPerInvoice", "TotalInvoices", "AvgLaborGM", "AvgPartsGM", "AvgEfficiency"]].corr()
ax = sns.heatmap(corr, annot=True, cmap="coolwarm", fmt=".2f", cbar_kws={'shrink': 0.97})

plt.title("Correlation Matrix: Customer-Level Numeric Features", fontsize=26, color='blue', fontweight='bold', pad=18)
plt.xticks(fontsize=15, color='white', rotation=35, ha='right')
plt.yticks(fontsize=15, color='white', rotation=45)
ax.set_facecolor('black')
for spine in ax.spines.values():
    spine.set_edgecolor('blue')
    spine.set_linewidth(2.1)

cbar = ax.collections[0].colorbar
cbar.ax.yaxis.set_tick_params(color='white', labelcolor='white')
cbar.outline.set_edgecolor('white')

plt.tight_layout()
plt.show()
```


Correlation Matrix Summary:


AvgTotalSalesPerInvoice	vs. TotalInvoices	: 0.00
AvgTotalSalesPerInvoice	vs. AvgLaborGM	: 0.40
AvgTotalSalesPerInvoice	vs. AvgPartsGM	: 0.18
AvgTotalSalesPerInvoice	vs. AvgEfficiency	: -0.07

TotalInvoices	vs. AvgTotalSalesPerInvoice:	0.00
TotalInvoices	vs. AvgLaborGM	: -0.02
TotalInvoices	vs. AvgPartsGM	: 0.06
TotalInvoices	vs. AvgEfficiency	: 0.08

AvgLaborGM	vs. AvgTotalSalesPerInvoice:	0.40
AvgLaborGM	vs. TotalInvoices	: -0.02
AvgLaborGM	vs. AvgPartsGM	: -0.06
AvgLaborGM	vs. AvgEfficiency	: 0.01

AvgPartsGM	vs. AvgTotalSalesPerInvoice:	0.18
AvgPartsGM	vs. TotalInvoices	: 0.06
AvgPartsGM	vs. AvgLaborGM	: -0.06
AvgPartsGM	vs. AvgEfficiency	: 0.04

AvgEfficiency	vs. AvgTotalSalesPerInvoice:	-0.07
AvgEfficiency	vs. TotalInvoices	: 0.08
AvgEfficiency	vs. AvgLaborGM	: 0.01
AvgEfficiency	vs. AvgPartsGM	: 0.04



```
In [21]: # Summary print outputs (keep as you have)
ti_median = customer_df["TotalInvoices"].median()
ti_mean = customer_df["TotalInvoices"].mean()
```

```

ti_min = customer_df["TotalInvoices"].min()
ti_max = customer_df["TotalInvoices"].max()
ti_25 = customer_df["TotalInvoices"].quantile(0.25)
ti_75 = customer_df["TotalInvoices"].quantile(0.75)

sales_median = customer_df["AvgTotalSalesPerInvoice"].median()
sales_mean = customer_df["AvgTotalSalesPerInvoice"].mean()
sales_min = customer_df["AvgTotalSalesPerInvoice"].min()
sales_max = customer_df["AvgTotalSalesPerInvoice"].max()
sales_25 = customer_df["AvgTotalSalesPerInvoice"].quantile(0.25)
sales_75 = customer_df["AvgTotalSalesPerInvoice"].quantile(0.75)

print("Total Invoices per Customer summary:")
print(f"  Median: {ti_median:.0f}")
print(f"  Mean:   {ti_mean:.1f}")
print(f"  Min:    {ti_min:.0f}")
print(f"  Max:    {ti_max:.0f}")
print(f"  25th percentile: {ti_25:.0f}")
print(f"  75th percentile: {ti_75:.0f}")
print()
print("Average Total Sales per Customer summary:")
print(f"  Median: ${sales_median:,.0f}")
print(f"  Mean:   ${sales_mean:,.0f}")
print(f"  Min:    ${sales_min:,.0f}")
print(f"  Max:    ${sales_max:,.0f}")
print(f"  25th percentile: ${sales_25:,.0f}")
print(f"  75th percentile: ${sales_75:,.0f}")

# Print Top 10 Customers by Total Invoices
print("Top 10 Customers by Total Invoices:")
top10_invoices = customer_df.nlargest(10, "TotalInvoices")
for idx, row in top10_invoices.iterrows():
    print(f"{idx+1:2d}. {row['CustName']:<25} | Invoices: {int(row['TotalInvoices']):5d} | Avg Total Sales: ${row['AvgTotalSalesPerInvoice']:.0f}")

print("\nTop 10 Customers by Average Total Sales per Invoice:")
top10_sales = customer_df.nlargest(10, "AvgTotalSalesPerInvoice")
for idx, row in top10_sales.iterrows():
    print(f"{idx+1:2d}. {row['CustName']:<25} | Avg Total Sales: ${row['AvgTotalSalesPerInvoice']:.0f} | Invoices: {int(row['TotalInvoices']):5d}")
print("\n" + "-"*70 + "\n")

# --- Annotate top customers ---
top3_sales = customer_df.nlargest(3, "AvgTotalSalesPerInvoice")
top3_invoices = customer_df.nlargest(3, "TotalInvoices")

plt.figure(figsize=(13, 8), facecolor='black')
ax = sns.scatterplot(
    x="TotalInvoices", y="AvgTotalSalesPerInvoice", data=customer_df,
    color='lime', edgecolor='crimson', s=110, alpha=0.75
)

plt.title("Total Invoices vs. Average Total Sales per Customer", fontsize=24, color='deepskyblue', fontweight='bold')
plt.xlabel("Total Invoices", fontsize=22, color='deepskyblue', fontweight='bold')
plt.ylabel("Average Total Sales per Customer", fontsize=22, color='deepskyblue', fontweight='bold')
plt.xticks(fontsize=14, color='white')
plt.yticks(fontsize=14, color='white')
ax.set_facecolor('black')
for spine in ax.spines.values():
    spine.set_edgecolor('cyan')
    spine.set_linewidth(2)
plt.tight_layout()

# Find top 3 by sales and by invoices
top3_sales = customer_df.nlargest(3, "AvgTotalSalesPerInvoice")
top3_invoices = customer_df.nlargest(3, "TotalInvoices")

# Manually set label offsets so they don't overlap (tweak as needed)
sales_offsets = [(120, -300), (120, 850), (120, 100)]
invoices_offsets = [(-450, 650), (150, -600), (-300, -1200)]

# Annotate Top 3 by Sales (yellow)
for i, (_, row) in enumerate(top3_sales.iterrows()):
    dx, dy = sales_offsets[i]
    plt.annotate(

```

```

f"#{i+1} Sales\n{row['CustName']}",
xy=(row["TotalInvoices"], row["AvgTotalSalesPerInvoice"]),
xytext=(row["TotalInvoices"]+dx, row["AvgTotalSalesPerInvoice"]+dy),
arrowprops=dict(facecolor='yellow', edgecolor='red', arrowstyle='->', linewidth=2.3),
fontsize=13, color='yellow', fontweight='bold',
bbox=dict(boxstyle='round', fc='black', ec='yellow', alpha=0.85)
)

# Annotate Top 3 by Total Invoices (magenta)
for i, (_, row) in enumerate(top3_invoices.iterrows()):
    dx, dy = invoices_offsets[i]
    plt.annotate(
        f"#{i+1} Invoices\n{row['CustName']}",
        xy=(row["TotalInvoices"], row["AvgTotalSalesPerInvoice"]),
        xytext=(row["TotalInvoices"]+dx, row["AvgTotalSalesPerInvoice"]+dy),
        arrowprops=dict(facecolor='magenta', edgecolor='red', arrowstyle='->', linewidth=2.3),
        fontsize=13, color='magenta', fontweight='bold',
        bbox=dict(boxstyle='round', fc='black', ec='magenta', alpha=0.85)
    )

plt.show()

```

Total Invoices per Customer summary:

Median: 28
 Mean: 117.6
 Min: 1
 Max: 2008
 25th percentile: 13
 75th percentile: 68

Average Total Sales per Customer summary:

Median: \$1,641
 Mean: \$1,729
 Min: \$86
 Max: \$5,563
 25th percentile: \$1,386
 75th percentile: \$1,997


Top 10 Customers by Total Invoices:

21. Paramount Transfer	Invoices: 2008 Avg Total Sales: \$1,927
121. Acme Inc	Invoices: 1747 Avg Total Sales: \$1,814
34. Paramount Logistics	Invoices: 1689 Avg Total Sales: \$2,007
85. Liberty Industries	Invoices: 1677 Avg Total Sales: \$1,867
142. Goldstar Transfer	Invoices: 1470 Avg Total Sales: \$1,662
151. Stonegate Partners	Invoices: 1375 Avg Total Sales: \$2,113
311. Unity Resources	Invoices: 1295 Avg Total Sales: \$1,797
305. Evergreen Associates	Invoices: 1212 Avg Total Sales: \$1,432
35. Aurora Inc	Invoices: 1193 Avg Total Sales: \$1,558
10. Evergreen LLC	Invoices: 1136 Avg Total Sales: \$1,186

Top 10 Customers by Average Total Sales per Invoice:

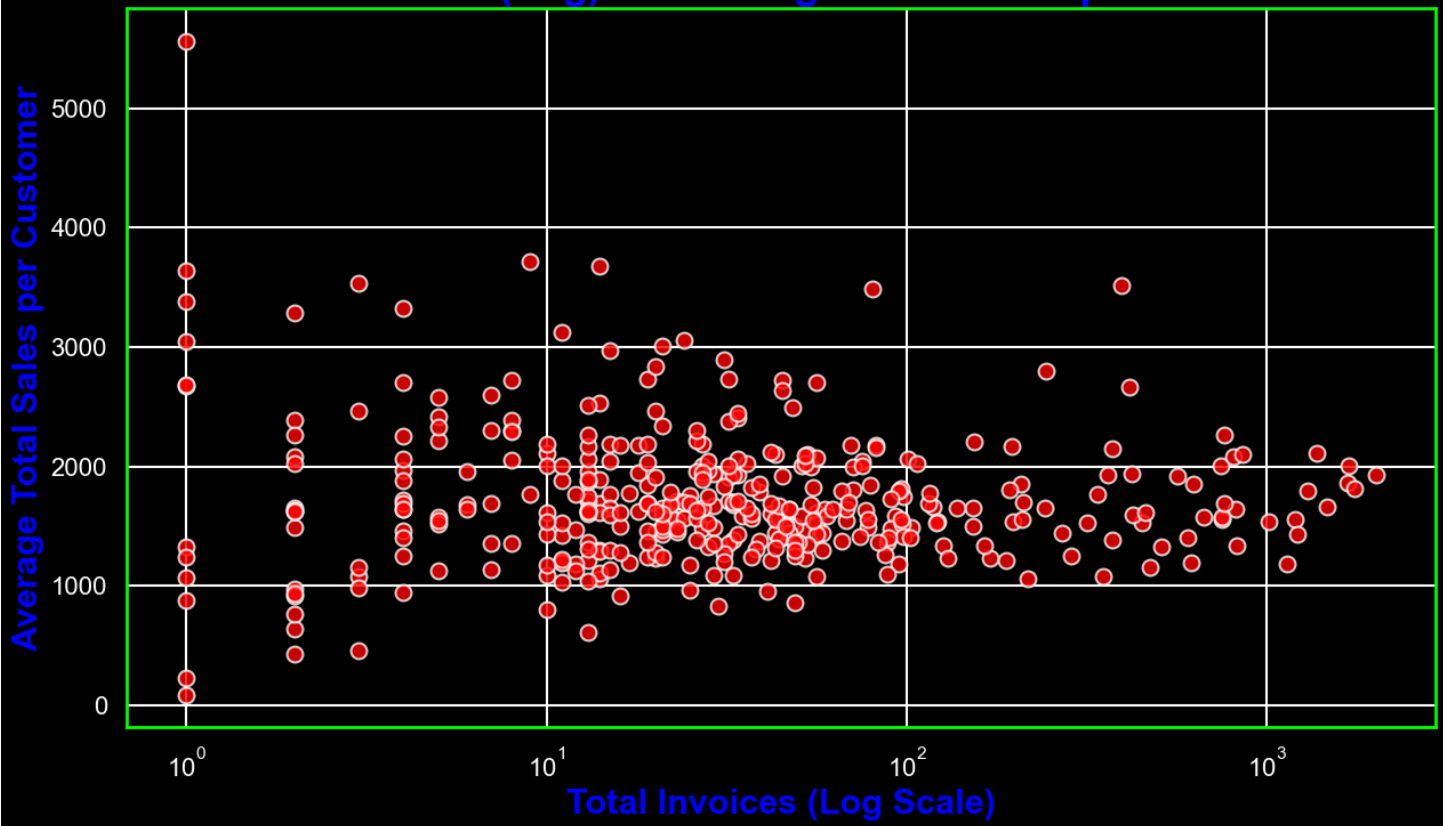
44. Pioneer Associates	Avg Total Sales: \$5,563 Invoices: 1
327. Keystone Tucking	Avg Total Sales: \$3,718 Invoices: 9
313. Pioneer Partners	Avg Total Sales: \$3,681 Invoices: 14
45. Pioneer Truck Lines	Avg Total Sales: \$3,637 Invoices: 1
221. Evergreen Transfer	Avg Total Sales: \$3,539 Invoices: 3
307. Sterling Systems	Avg Total Sales: \$3,520 Invoices: 395
149. Synergy Solutions	Avg Total Sales: \$3,492 Invoices: 80
192. Sterling Corp	Avg Total Sales: \$3,383 Invoices: 1
37. Evergreen Transfer	Avg Total Sales: \$3,329 Invoices: 4
65. Keystone Global	Avg Total Sales: \$3,287 Invoices: 2

Total Invoices vs. Average Total Sales per Customer



```
In [22]: plt.figure(figsize=(13, 8), facecolor='black')
ax = plt.gca()
plt.scatter(customer_df["TotalInvoices"], customer_df["AvgTotalSalesPerInvoice"],
            s=100, c='red', edgecolors='white', alpha=0.8)
plt.xscale('log')
plt.title("Total Invoices (Log) vs. Average Total Sales per Customer", fontsize=26, color='blue', fontweight='bold')
plt.xlabel("Total Invoices (Log Scale)", fontsize=22, color='blue', fontweight='bold')
plt.ylabel("Average Total Sales per Customer", fontsize=22, color='blue', fontweight='bold')
plt.xticks(fontsize=16, color='white')
plt.yticks(fontsize=16, color='white')
ax.set_facecolor('black')
for spine in ax.spines.values():
    spine.set_edgecolor('lime')
    spine.set_linewidth(2)
plt.tight_layout()
plt.show()
```

Total Invoices (Log) vs. Average Total Sales per Customer



```
In [23]: # --- SERVICE ONLY: Group by customer for service invoices ---
service_mask = df["R0type"].isin(["RESALE", "TRUCK", "TRAILER"])
service_group = df[service_mask].groupby(["CustNo", "CustName"]).agg(
    TotalServiceInvoices=("InvoiceNo", "count"),
    AvgLaborSalesPerInvoice=("LaborBilled$", "mean"),
).reset_index()

# --- Print summaries ---
labor_median = service_group["AvgLaborSalesPerInvoice"].median()
labor_mean = service_group["AvgLaborSalesPerInvoice"].mean()
labor_min = service_group["AvgLaborSalesPerInvoice"].min()
labor_max = service_group["AvgLaborSalesPerInvoice"].max()
labor_25 = service_group["AvgLaborSalesPerInvoice"].quantile(0.25)
labor_75 = service_group["AvgLaborSalesPerInvoice"].quantile(0.75)

print("Average Labor Sales per Service Customer summary:")
print(f"  Median: ${labor_median:,.0f}")
print(f"  Mean:   ${labor_mean:,.0f}")
print(f"  Min:    ${labor_min:,.0f}")
print(f"  Max:    ${labor_max:,.0f}")
print(f"  25th percentile: ${labor_25:,.0f}")
print(f"  75th percentile: ${labor_75:,.0f}")

print("\nTop 10 Customers by Total Service Invoices:")
top10_service_invoices = service_group.nlargest(10, "TotalServiceInvoices")
for idx, row in top10_service_invoices.iterrows():
    print(f"{idx+1:2d}. {row['CustName'][:25]} | Service Invoices: {int(row['TotalServiceInvoices']):5d} | Avg Labor Sales: ${row['AvgLaborSalesPerInvoice']:,.0f}")

print("\nTop 10 Customers by Avg Labor Sales per Service Invoice:")
top10_labor_sales = service_group.nlargest(10, "AvgLaborSalesPerInvoice")
for idx, row in top10_labor_sales.iterrows():
    print(f"{idx+1:2d}. {row['CustName'][:25]} | Avg Labor Sales: ${row['AvgLaborSalesPerInvoice']:,.0f} | Service Invoices: {int(row['TotalServiceInvoices']):5d}")
print("\n" + "-"*70 + "\n")

# --- Annotate top 3 by labor sales and top 3 by service invoices ---
top3_labor = service_group.nlargest(3, "AvgLaborSalesPerInvoice")
top3_servinv = service_group.nlargest(3, "TotalServiceInvoices")

plt.figure(figsize=(13, 8), facecolor='black')
ax = sns.scatterplot(
    x="TotalServiceInvoices", y="AvgLaborSalesPerInvoice", data=service_group,
```

```

        color='springgreen', edgecolor='red', s=180, alpha=0.80
    )
plt.title("Total Service Invoices vs. Average Labor Sales per Customer", fontsize=26, color='springgreen', fontweight='bold')
plt.xlabel("Total Service Invoices", fontsize=22, color='springgreen', fontweight='bold')
plt.ylabel("Average Labor Sales per Service Customer", fontsize=22, color='springgreen', fontweight='bold')
plt.xticks(fontsize=16, color='white')
plt.yticks(fontsize=16, color='white')
ax.set_facecolor('black')
for spine in ax.spines.values():
    spine.set_edgecolor('cyan')
    spine.set_linewidth(2)
plt.tight_layout()

sales_offsets = [(100, -200), (60, -800), (60, -1200)]
invoices_offsets = [(-175, 650), (25, -750), (-150, -1400)]

for i, (_, row) in enumerate(top3_labor.iterrows()):
    dx, dy = sales_offsets[i]
    plt.annotate(
        f"#{i+1} Labor Sales\n{row['CustName']}",
        xy=(row["TotalServiceInvoices"], row["AvgLaborSalesPerInvoice"]),
        xytext=(row["TotalServiceInvoices"]+dx, row["AvgLaborSalesPerInvoice"]+dy),
        arrowprops=dict(facecolor='yellow', edgecolor='red', arrowstyle='->', linewidth=2.2),
        fontsize=13, color='yellow', fontweight='bold',
        bbox=dict(boxstyle='round', fc='black', ec='yellow', alpha=0.85)
    )

for i, (_, row) in enumerate(top3_servinv.iterrows()):
    dx, dy = invoices_offsets[i]
    plt.annotate(
        f"#{i+1} Service Invoices\n{row['CustName']}",
        xy=(row["TotalServiceInvoices"], row["AvgLaborSalesPerInvoice"]),
        xytext=(row["TotalServiceInvoices"]+dx, row["AvgLaborSalesPerInvoice"]+dy),
        arrowprops=dict(facecolor='magenta', edgecolor='red', arrowstyle='->', linewidth=2.2),
        fontsize=13, color='magenta', fontweight='bold',
        bbox=dict(boxstyle='round', fc='black', ec='magenta', alpha=0.85)
    )

plt.show()

```

Average Labor Sales per Service Customer summary:


Median: \$1,920
Mean: \$2,057
Min: \$729
Max: \$5,487
25th percentile: \$1,621
75th percentile: \$2,341

Top 10 Customers by Total Service Invoices:

19. Paramount Transfer	Service Invoices: 1176	Avg Labor Sales: \$2,028
31. Paramount Logistics	Service Invoices: 923	Avg Labor Sales: \$1,872
82. Liberty Industries	Service Invoices: 875	Avg Labor Sales: \$2,088
118. Acme Inc	Service Invoices: 847	Avg Labor Sales: \$1,830
146. Stonegate Partners	Service Invoices: 726	Avg Labor Sales: \$2,366
137. Goldstar Transfer	Service Invoices: 715	Avg Labor Sales: \$1,635
32. Aurora Inc	Service Invoices: 588	Avg Labor Sales: \$1,444
301. Unity Resources	Service Invoices: 569	Avg Labor Sales: \$2,306
295. Evergreen Associates	Service Invoices: 557	Avg Labor Sales: \$1,538
353. Acme Systems	Service Invoices: 463	Avg Labor Sales: \$2,273

Top 10 Customers by Avg Labor Sales per Service Invoice:

303. Pioneer Partners	Avg Labor Sales: \$5,487	Service Invoices: 7
144. Synergy Solutions	Avg Labor Sales: \$5,406	Service Invoices: 39
95. Acme Logistics	Avg Labor Sales: \$5,165	Service Invoices: 1
41. Pioneer Associates	Avg Labor Sales: \$4,464	Service Invoices: 1
26. Stonegate Truck Lines	Avg Labor Sales: \$4,382	Service Invoices: 120
317. Keystone Tucking	Avg Labor Sales: \$4,235	Service Invoices: 6
297. Sterling Systems	Avg Labor Sales: \$4,234	Service Invoices: 269
36. Paramount Global	Avg Labor Sales: \$3,911	Service Invoices: 7
240. Synergy Transfer	Avg Labor Sales: \$3,864	Service Invoices: 215
200. Pioneer Partners	Avg Labor Sales: \$3,862	Service Invoices: 22



```
In [24]: # --- Calculate for ALL invoices: Group by customer ---
parts_group = df.groupby(["CustNo", "CustName"]).agg(
    TotalInvoices=("InvoiceNo", "count"),
    AvgPartsSalesPerInvoice=("PartsSales$", "mean"),
).reset_index()

# Merge for annotation (optional)
```



```

customer_df = customer_df.merge(parts_group, on=["CustNo", "CustName"], suffixes=('_', '_all'), how='left')

# Print summaries
parts_median = parts_group["AvgPartsSalesPerInvoice"].median()
parts_mean = parts_group["AvgPartsSalesPerInvoice"].mean()
parts_min = parts_group["AvgPartsSalesPerInvoice"].min()
parts_max = parts_group["AvgPartsSalesPerInvoice"].max()
parts_25 = parts_group["AvgPartsSalesPerInvoice"].quantile(0.25)
parts_75 = parts_group["AvgPartsSalesPerInvoice"].quantile(0.75)

print("Average Parts Sales per Customer summary:")
print(f"  Median:  ${parts_median:,.0f}")
print(f"  Mean:    ${parts_mean:,.0f}")
print(f"  Min:     ${parts_min:,.0f}")
print(f"  Max:     ${parts_max:,.0f}")
print(f"  25th percentile: ${parts_25:,.0f}")
print(f"  75th percentile: ${parts_75:,.0f}")

print("\nTop 10 Customers by Total Invoices:")
top10_inv = parts_group.nlargest(10, "TotalInvoices")
for idx, row in top10_inv.iterrows():
    print(f"{idx+1:2d}. {row['CustName'][:25]} | Invoices: {int(row['TotalInvoices']):5d} | Avg Parts Sales: ${row['AvgPartsSalesPerInvoice']:.0f}")

print("\nTop 10 Customers by Avg Parts Sales per Invoice:")
top10_part_sales = parts_group.nlargest(10, "AvgPartsSalesPerInvoice")
for idx, row in top10_part_sales.iterrows():
    print(f"{idx+1:2d}. {row['CustName'][:25]} | Avg Parts Sales: ${row['AvgPartsSalesPerInvoice']:.0f} | Invoices: {int(row['TotalInvoices']):5d}")
print("\n" + "-"*70 + "\n")

# Top 3 by Avg Parts Sales and by Invoices
top3_parts = parts_group.nlargest(3, "AvgPartsSalesPerInvoice")
top3_inv = parts_group.nlargest(3, "TotalInvoices")

plt.figure(figsize=(13, 8), facecolor='black')
ax = sns.scatterplot(
    x="TotalInvoices", y="AvgPartsSalesPerInvoice", data=parts_group,
    color='orange', edgecolor='lime', s=110, alpha=0.80
)
plt.title("Total Invoices vs. Average Parts Sales per Customer", fontsize=26, color='orange', fontweight='bold', pad=10)
plt.xlabel("Total Invoices", fontsize=22, color='orange', fontweight='bold')
plt.ylabel("Average Parts Sales per Customer", fontsize=22, color='orange', fontweight='bold')
plt.xticks(fontsize=16, color='white')
plt.yticks(fontsize=16, color='white')
ax.set_facecolor('black')
for spine in ax.spines.values():
    spine.set_edgecolor('lime')
    spine.set_linewidth(2)
plt.tight_layout()

sales_offsets = [(250, -100), (300, 50), (300, -100)]
invoices_offsets = [(-300, 350), (100, -800), (-200, -700)]

for i, (_, row) in enumerate(top3_parts.iterrows()):
    dx, dy = sales_offsets[i]
    plt.annotate(
        f"#{i+1} Parts Sales\n{row['CustName']}",
        xy=(row["TotalInvoices"], row["AvgPartsSalesPerInvoice"]),
        xytext=(row["TotalInvoices"]+dx, row["AvgPartsSalesPerInvoice"]+dy),
        arrowprops=dict(facecolor='yellow', edgecolor='red', arrowstyle='->', linewidth=2.2),
        fontsize=13, color='yellow', fontweight='bold',
        bbox=dict(boxstyle='round', fc='black', ec='yellow', alpha=0.85)
    )

for i, (_, row) in enumerate(top3_inv.iterrows()):
    dx, dy = invoices_offsets[i]
    plt.annotate(
        f"#{i+1} Invoices\n{row['CustName']}",
        xy=(row["TotalInvoices"], row["AvgPartsSalesPerInvoice"]),
        xytext=(row["TotalInvoices"]+dx, row["AvgPartsSalesPerInvoice"]+dy),
        arrowprops=dict(facecolor='magenta', edgecolor='red', arrowstyle='->', linewidth=2.2),
        fontsize=13, color='magenta', fontweight='bold',
        bbox=dict(boxstyle='round', fc='black', ec='magenta', alpha=0.85)
    )

```



```
plt.show()
```

Average Parts Sales per Customer summary:


Median: \$802
Mean: \$804
Min: \$86
Max: \$1,750
25th percentile: \$697
75th percentile: \$893

Top 10 Customers by Total Invoices:

21. Paramount Transfer	Invoices: 2008 Avg Parts Sales: \$739
121. Acme Inc	Invoices: 1747 Avg Parts Sales: \$927
34. Paramount Logistics	Invoices: 1689 Avg Parts Sales: \$984
85. Liberty Industries	Invoices: 1677 Avg Parts Sales: \$777
142. Goldstar Transfer	Invoices: 1470 Avg Parts Sales: \$867
151. Stonegate Partners	Invoices: 1375 Avg Parts Sales: \$864
311. Unity Resources	Invoices: 1295 Avg Parts Sales: \$784
305. Evergreen Associates	Invoices: 1212 Avg Parts Sales: \$725
35. Aurora Inc	Invoices: 1193 Avg Parts Sales: \$846
10. Evergreen LLC	Invoices: 1136 Avg Parts Sales: \$1,005

Top 10 Customers by Avg Parts Sales per Invoice:

275. Goldstar Solutions	Avg Parts Sales: \$1,750 Invoices: 1
159. Unity Inc	Avg Parts Sales: \$1,327 Invoices: 1
82. Oakwood Truck Lines	Avg Parts Sales: \$1,286 Invoices: 8
205. Sunset Solutions	Avg Parts Sales: \$1,254 Invoices: 3
381. Evergreen Tucking	Avg Parts Sales: \$1,228 Invoices: 5
283. Liberty Inc	Avg Parts Sales: \$1,203 Invoices: 8
207. Summit Tucking	Avg Parts Sales: \$1,183 Invoices: 5
37. Evergreen Transfer	Avg Parts Sales: \$1,164 Invoices: 4
141. Goldstar Global	Avg Parts Sales: \$1,161 Invoices: 15
16. Stonegate Group	Avg Parts Sales: \$1,158 Invoices: 2



```
In [25]: # ----- Print summary stats -----
eff_median = customer_df["AvgEfficiency"].median()
eff_mean = customer_df["AvgEfficiency"].mean()
eff_min = customer_df["AvgEfficiency"].min()
eff_max = customer_df["AvgEfficiency"].max()
```

```

eff_25 = customer_df["AvgEfficiency"].quantile(0.25)
eff_75 = customer_df["AvgEfficiency"].quantile(0.75)

sales_median = customer_df["AvgTotalSalesPerInvoice"].median()
sales_mean = customer_df["AvgTotalSalesPerInvoice"].mean()
sales_min = customer_df["AvgTotalSalesPerInvoice"].min()
sales_max = customer_df["AvgTotalSalesPerInvoice"].max()
sales_25 = customer_df["AvgTotalSalesPerInvoice"].quantile(0.25)
sales_75 = customer_df["AvgTotalSalesPerInvoice"].quantile(0.75)

print("Average Efficiency per Customer summary:")
print(f"  Median: {eff_median:.2f}")
print(f"  Mean:   {eff_mean:.2f}")
print(f"  Min:    {eff_min:.2f}")
print(f"  Max:    {eff_max:.2f}")
print(f"  25th percentile: {eff_25:.2f}")
print(f"  75th percentile: {eff_75:.2f}\n")

print("Average Total Sales per Customer summary:")
print(f"  Median: ${sales_median:,.0f}")
print(f"  Mean:   ${sales_mean:,.0f}")
print(f"  Min:    ${sales_min:,.0f}")
print(f"  Max:    ${sales_max:,.0f}")
print(f"  25th percentile: ${sales_25:,.0f}")
print(f"  75th percentile: ${sales_75:,.0f}")

# Top 10 customers by efficiency and by average sales
print("\nTop 10 Customers by Average Efficiency:")
top10_eff = customer_df.nlargest(10, "AvgEfficiency")
for idx, row in top10_eff.iterrows():
    print(f"{idx+1:2d}. {row['CustName']:<25} | Efficiency: {row['AvgEfficiency']:.2f} | Avg Total Sales: ${row['AvgTotalSalesPerInvoice']:.0f}")

print("\nTop 10 Customers by Average Total Sales per Invoice:")
top10_sales = customer_df.nlargest(10, "AvgTotalSalesPerInvoice")
for idx, row in top10_sales.iterrows():
    print(f"{idx+1:2d}. {row['CustName']:<25} | Avg Total Sales: ${row['AvgTotalSalesPerInvoice']:.0f} | Efficiency: {row['AvgEfficiency']:.2f}")
print("\n" + "-"*70 + "\n")

# --- Annotate top customers ---
top3_sales = customer_df.nlargest(3, "AvgTotalSalesPerInvoice")
top3_eff = customer_df.nlargest(3, "AvgEfficiency")

plt.figure(figsize=(13, 8), facecolor='black')
ax = sns.scatterplot(
    x="AvgEfficiency", y="AvgTotalSalesPerInvoice", data=customer_df,
    color='orange', edgecolor='crimson', s=110, alpha=0.77
)

plt.title("Average Efficiency vs. Average Total Sales per Customer", fontsize=26, color='deepskyblue', fontweight='bold')
plt.xlabel("Average Efficiency", fontsize=22, color='deepskyblue', fontweight='bold')
plt.ylabel("Average Total Sales per Customer", fontsize=22, color='deepskyblue', fontweight='bold')
plt.xticks(fontsize=14, color='white')
plt.yticks(fontsize=14, color='white')
ax.set_facecolor('black')
for spine in ax.spines.values():
    spine.set_edgecolor('cyan')
    spine.set_linewidth(2)
plt.tight_layout()

# Label offsets (tweak for your data layout)
sales_offsets = [(-0.1, -900), (-0.05, 625), (-0.1, 300)]
eff_offsets = [(-0.14, 950), (-0.08, -1200), (-0.25, 700)]

# Annotate Top 3 by Sales (yellow)
for i, (_, row) in enumerate(top3_sales.iterrows()):
    dx, dy = sales_offsets[i]
    plt.annotate(
        f"#{i+1} Sales\n{row['CustName']}",
        xy=(row["AvgEfficiency"], row["AvgTotalSalesPerInvoice"]),
        xytext=(row["AvgEfficiency"]+dx, row["AvgTotalSalesPerInvoice"]+dy),
        arrowprops=dict(facecolor='yellow', edgecolor='red', arrowstyle='->', linewidth=2.3),
        fontsize=13, color='yellow', fontweight='bold',
        bbox=dict(boxstyle='round', fc='black', ec='yellow', alpha=0.85)
    )

```

```

)
# Annotate Top 3 by Efficiency (lime)
for i, (_, row) in enumerate(top3_eff.iterrows()):
    dx, dy = eff_offsets[i]
    plt.annotate(
        f"#{i+1} Efficiency\n{row['CustName']}",
        xy=(row["AvgEfficiency"], row["AvgTotalSalesPerInvoice"]),
        xytext=(row["AvgEfficiency"]+dx, row["AvgTotalSalesPerInvoice"]+dy),
        arrowprops=dict(facecolor='lime', edgecolor='red', arrowstyle='->', linewidth=2.3),
        fontsize=13, color='lime', fontweight='bold',
        bbox=dict(boxstyle='round', fc='black', ec='lime', alpha=0.85)
    )
plt.show()

```

Average Efficiency per Customer summary:

```

Median: 1.24
Mean:    1.24
Min:     0.60
Max:     1.84
25th percentile: 1.11
75th percentile: 1.38

```

Average Total Sales per Customer summary:

```

Median: $1,641
Mean:    $1,729
Min:     $86
Max:     $5,563
25th percentile: $1,386
75th percentile: $1,997

```

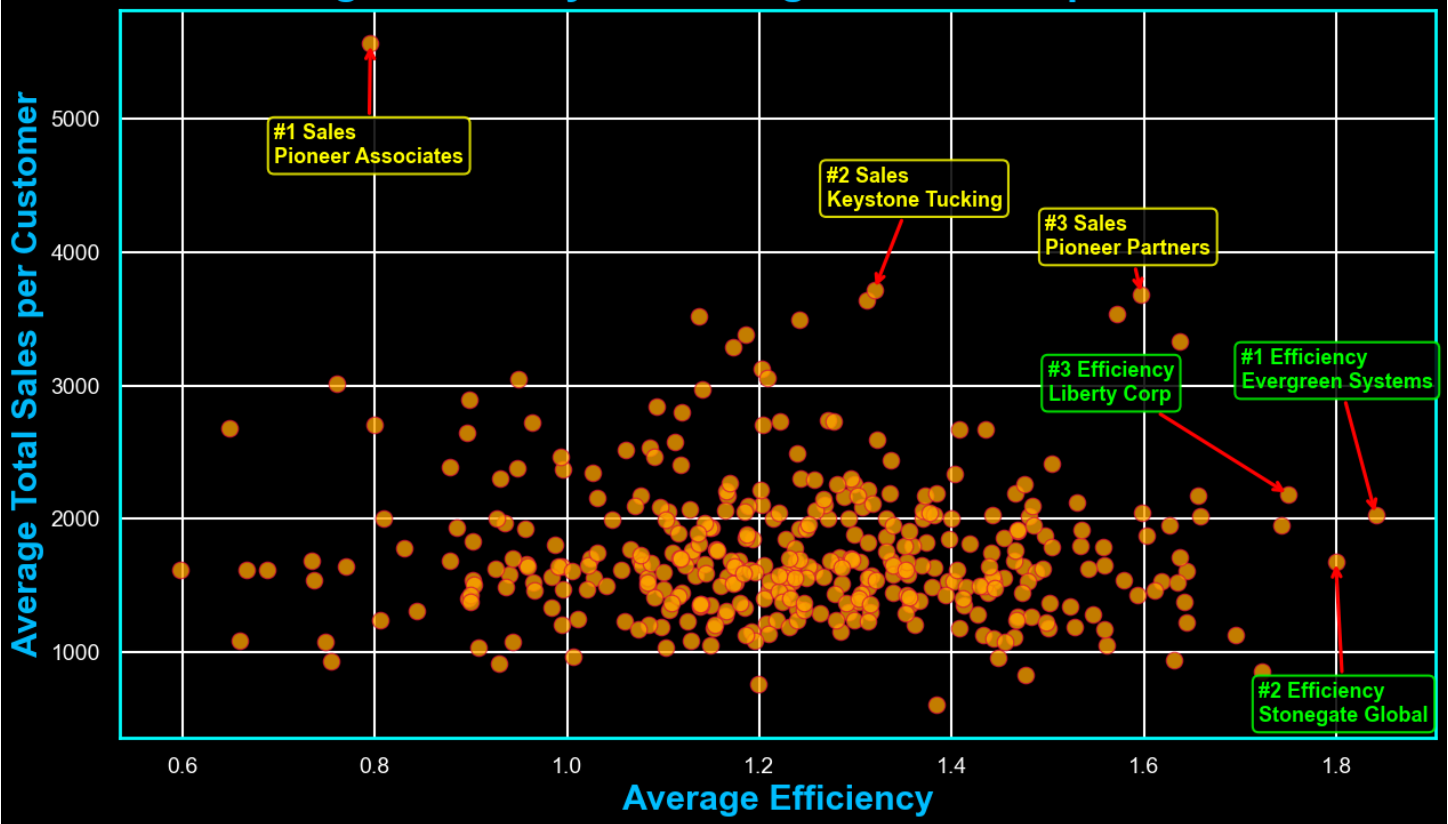
Top 10 Customers by Average Efficiency:

125. Evergreen Systems	Efficiency: 1.84	Avg Total Sales: \$2,029
153. Stonegate Global	Efficiency: 1.80	Avg Total Sales: \$1,683
111. Liberty Corp	Efficiency: 1.75	Avg Total Sales: \$2,184
18. Keystone Industries	Efficiency: 1.74	Avg Total Sales: \$1,956
317. Evergreen LLC	Efficiency: 1.72	Avg Total Sales: \$857
301. Keystone Inc	Efficiency: 1.70	Avg Total Sales: \$1,135
226. Keystone Associates	Efficiency: 1.66	Avg Total Sales: \$2,020
69. Liberty Associates	Efficiency: 1.66	Avg Total Sales: \$2,176
181. Unity LLC	Efficiency: 1.64	Avg Total Sales: \$1,229
51. Evergreen Logistics	Efficiency: 1.64	Avg Total Sales: \$1,608

Top 10 Customers by Average Total Sales per Invoice:

44. Pioneer Associates	Avg Total Sales: \$5,563	Efficiency: 0.80
327. Keystone Tucking	Avg Total Sales: \$3,718	Efficiency: 1.32
313. Pioneer Partners	Avg Total Sales: \$3,681	Efficiency: 1.60
45. Pioneer Truck Lines	Avg Total Sales: \$3,637	Efficiency: 1.31
221. Evergreen Transfer	Avg Total Sales: \$3,539	Efficiency: 1.57
307. Sterling Systems	Avg Total Sales: \$3,520	Efficiency: 1.14
149. Synergy Solutions	Avg Total Sales: \$3,492	Efficiency: 1.24
192. Sterling Corp	Avg Total Sales: \$3,383	Efficiency: 1.19
37. Evergreen Transfer	Avg Total Sales: \$3,329	Efficiency: 1.64
65. Keystone Global	Avg Total Sales: \$3,287	Efficiency: 1.17

Average Efficiency vs. Average Total Sales per Customer



```
In [26]: plt.figure(figsize=(13, 7), facecolor='black')

# Print summary: Average total sales by Most Common R0type
rotype_means = customer_df.groupby("MostCommonR0type")["AvgTotalSalesPerInvoice"].agg(['mean', 'count']).sort_values(

print("Average Total Sales per Customer by Most Common R0type:")
for idx, row in rotype_means.iterrows():
    print(f" {idx:12s} | Avg Sale: ${row['mean']:, .0f} | Customers: {int(row['count'])}")
print()

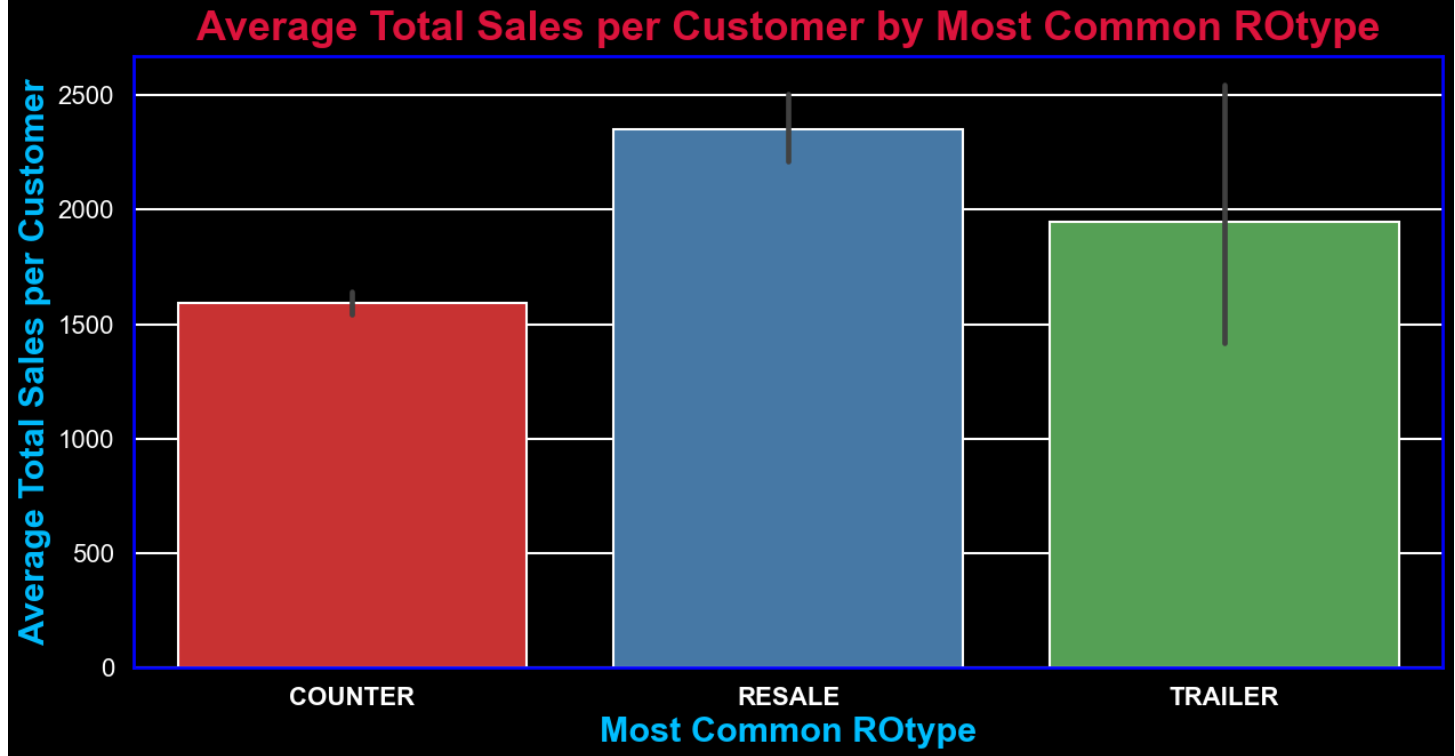
palette = sns.color_palette("bright", n_colors=customer_df["MostCommonR0type"].nunique())

ax = sns.barplot(
    x="MostCommonR0type",
    y="AvgTotalSalesPerInvoice",
    data=customer_df,
    hue="MostCommonR0type",
    palette="Set1",
    legend=False
)

plt.title("Average Total Sales per Customer by Most Common R0type", fontsize=26, color='crimson', fontweight='bold',
plt.xlabel("Most Common R0type", fontsize=22, color='deepskyblue', fontweight='bold')
plt.ylabel("Average Total Sales per Customer", fontsize=22, color='deepskyblue', fontweight='bold')
plt.xticks(fontsize=16, color='white', fontweight='bold')
plt.yticks(fontsize=16, color='white')
ax.set_facecolor('black')
for spine in ax.spines.values():
    spine.set_edgecolor('blue')
    spine.set_linewidth(2)
plt.tight_layout()
plt.show()
```

Average Total Sales per Customer by Most Common R0type:

RESALE	Avg Sale: \$2,352	Customers: 67
TRAILER	Avg Sale: \$1,944	Customers: 5
COUNTER	Avg Sale: \$1,593	Customers: 315



Statistical Test: Difference Between RESALE and COUNTER Customers

We use a t-test to see if average total sales differ significantly between RESALE and COUNTER dominant customers.

```
In [27]: from scipy.stats import ttest_ind

# Filter customers by most common ROtype
resale_sales = customer_df[customer_df["MostCommonROtype"] == "RESALE"]["AvgTotalSalesPerInvoice"]
counter_sales = customer_df[customer_df["MostCommonROtype"] == "COUNTER"]["AvgTotalSalesPerInvoice"]

# Perform two-sample t-test
t_stat, p_value = ttest_ind(resale_sales, counter_sales, nan_policy='omit')
print("T-statistic:", t_stat)
print("P-value:", p_value)
```

```
T-statistic: 11.407577033806557
P-value: 4.0199041135806196e-26
```

Prepare Data for Modeling

We select numeric and categorical features, encode categorical fields, and fill any missing values before modeling.

```
In [28]: # Prepare features for modeling (drop ID columns, encode categoricals)
X = customer_df.drop(columns=["CustNo", "CustName", "AvgTotalSalesPerInvoice"])
X = pd.get_dummies(X, columns=["MostCommonROtype", "MostCommonDept", "MostCommonLocation"], drop_first=True)

# Fill any missing values with zero (safe for dummy/agg columns)
X = X.fillna(0)

y = customer_df["AvgTotalSalesPerInvoice"]
X = X.loc[:, ~X.columns.duplicated()]

# Print statements to confirm prep
print("Feature matrix X shape:", X.shape)
print("Target vector y shape:", y.shape)
print("First 5 rows of feature matrix X:")
print(X.head())
print("\nFeature columns used in modeling:")
print(list(X.columns))
```

Feature matrix X shape: (387, 17)

Target vector y shape: (387,)

First 5 rows of feature matrix X:

	TotalInvoices	AvgLaborGM	AvgPartsGM	AvgEfficiency	TotalInvoices_all	\
0	742	0.658874	0.345703	1.370037	742	
1	23	0.574704	0.223554	1.024120	23	
2	82	0.611540	0.232640	1.076672	82	
3	13	0.592245	0.323652	1.501356	13	
4	11	0.756308	0.335894	1.298439	11	

	AvgPartsSalesPerInvoice	MostCommonROtype_RESALE	MostCommonROtype_TRAILER	\
0	925.104474	False	False	
1	675.424783	False	False	
2	825.246585	True	False	
3	724.214615	False	False	
4	731.129091	False	False	

	MostCommonDept_20	MostCommonDept_30	MostCommonDept_40	MostCommonDept_50	\
0	True	False	False	False	
1	False	False	False	False	
2	False	False	True	False	
3	True	False	False	False	
4	False	False	True	False	

	MostCommonLocation_Chicago	MostCommonLocation_Dallas	\
0	False	False	
1	False	False	
2	False	False	
3	False	True	
4	False	False	

	MostCommonLocation_Green Bay	MostCommonLocation_Los Angeles	\
0	False	True	
1	True	False	
2	False	False	
3	False	False	
4	True	False	

	MostCommonLocation_New York
0	False
1	False
2	False
3	False
4	False

Feature columns used in modeling:

```
['TotalInvoices', 'AvgLaborGM', 'AvgPartsGM', 'AvgEfficiency', 'TotalInvoices_all', 'AvgPartsSalesPerInvoice', 'MostCommonROtype_RESALE', 'MostCommonROtype_TRAILER', 'MostCommonDept_20', 'MostCommonDept_30', 'MostCommonDept_40', 'MostCommonDept_50', 'MostCommonLocation_Chicago', 'MostCommonLocation_Dallas', 'MostCommonLocation_Green Bay', 'MostCommonLocation_Los Angeles', 'MostCommonLocation_New York']
```

Train/Test Split

We split the customer data into training and test sets.

```
In [29]: from sklearn.model_selection import train_test_split

# Split data into training and test sets (80% train, 20% test)
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42
)

print("Train/Test Split Results:")
print(f" X_train shape: {X_train.shape}")
print(f" X_test shape: {X_test.shape}")
print(f" y_train shape: {y_train.shape}")
print(f" y_test shape: {y_test.shape}")
print(f" Training set percent: {100*len(X_train)/len(X):.1f}%")
print(f" Test set percent: {100*len(X_test)/len(X):.1f}%")
```

Train/Test Split Results:
X_train shape: (309, 17)
X_test shape: (78, 17)
y_train shape: (309,)
y_test shape: (78,)
Training set percent: 79.8%
Test set percent: 20.2%

Model 1: Linear Regression

We begin with a linear regression baseline and report its metrics.

```
In [30]: # Import model and metrics for linear regression
from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score, mean_absolute_error, mean_squared_error

# If available, import the new RMSE function for future compatibility
try:
    from sklearn.metrics import root_mean_squared_error
    use_new_rmse = True
except ImportError:
    use_new_rmse = False

# Train a linear regression model on the training data
lr = LinearRegression()
lr.fit(X_train, y_train)

# Make predictions on the test set
y_pred_lr = lr.predict(X_test)

# Evaluate model performance with common regression metrics:
# R2: How well the model explains variance (1 = perfect, 0 = average)
# MAE: Average prediction error in the same units as the target (dollars)
# RMSE: Root mean squared error; larger errors penalized more heavily.
# Note: As of scikit-learn 1.4, the preferred way to compute RMSE is with root_mean_squared_error().
# Older code often uses mean_squared_error(..., squared=False), which is being deprecated.

print("----- Linear Regression Model Performance -----")
print(f" Features used: {X_train.shape[1]}")
print(f" Test set size: {len(X_test)} customers\n")
print(f" R2 (Test):      {r2_score(y_test, y_pred_lr):.6f}")
print(f" MAE (Test):      ${mean_absolute_error(y_test, y_pred_lr):.2f}")
if use_new_rmse:
    print(f" RMSE (Test):      ${root_mean_squared_error(y_test, y_pred_lr):.2f}")
else:
    print(f" RMSE (Test):      ${mean_squared_error(y_test, y_pred_lr, squared=False):.2f}")
print("\n-----")
```

```
----- Linear Regression Model Performance -----
Features used: 17
Test set size: 78 customers

R2 (Test):      0.567384
MAE (Test):      $236.47
RMSE (Test):      $302.15

-----
```

Model 2: Random Forest Regressor

Next, we use a random forest to capture nonlinear effects and feature interactions.

```
In [31]: # Import model and metrics for random forest regression
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import r2_score, mean_absolute_error, mean_squared_error

try:
    from sklearn.metrics import root_mean_squared_error
    use_new_rmse = True
except ImportError:
```

```

use_new_rmse = False

# Train the random forest model
rf = RandomForestRegressor(n_estimators=100, random_state=42)
rf.fit(X_train, y_train)

# Predict on test set
y_pred_rf = rf.predict(X_test)

# Prepare output
n_features = X_train.shape[1]
n_test = len(y_test)
r2 = r2_score(y_test, y_pred_rf)
mae = mean_absolute_error(y_test, y_pred_rf)
rmse = root_mean_squared_error(y_test, y_pred_rf) if use_new_rmse else mean_squared_error(y_test, y_pred_rf, squared=

print("----- Random Forest Model Performance -----")
print(f"  Features used: {n_features}")
print(f"  Test set size: {n_test} customers\n")
print(f"  R² (Test):      {r2:.6f}")
print(f"  MAE (Test):     ${mae:,.2f}")
print(f"  RMSE (Test):    ${rmse:,.2f}\n")
print("-----")

```

```

----- Random Forest Model Performance -----
Features used: 17
Test set size: 78 customers

R² (Test):      0.525416
MAE (Test):     $250.68
RMSE (Test):    $316.47

```

In [32]: # Model 3: XGBoost Regressor

```

try:
    from xgboost import XGBRegressor
    from sklearn.metrics import r2_score, mean_absolute_error, mean_squared_error

    xgb = XGBRegressor(n_estimators=100, random_state=42)
    xgb.fit(X_train, y_train)
    y_pred_xgb = xgb.predict(X_test)

    try:
        from sklearn.metrics import root_mean_squared_error
        rmse = root_mean_squared_error(y_test, y_pred_xgb)
    except ImportError:
        rmse = mean_squared_error(y_test, y_pred_xgb, squared=False)

    n_features = X_train.shape[1]
    n_test = len(y_test)
    r2 = r2_score(y_test, y_pred_xgb)
    mae = mean_absolute_error(y_test, y_pred_xgb)

    print("----- XGBoost Model Performance -----")
    print(f"  Features used: {n_features}")
    print(f"  Test set size: {n_test} customers\n")
    print(f"  R² (Test):      {r2:.6f}")
    print(f"  MAE (Test):     ${mae:,.2f}")
    print(f"  RMSE (Test):    ${rmse:,.2f}\n")
    print("-----")
except ImportError:
    print("XGBoost is not installed. Please install xgboost using pip or conda and restart the kernel.")

```


----- XGBoost Model Performance -----

Features used: 17

Test set size: 78 customers

R² (Test): 0.426515

MAE (Test): \$269.12

RMSE (Test): \$347.88

Model Evaluation: Visualizations

We visualize predicted vs actual and residuals for the best-performing model.

```
In [33]: # Use predictions from your best model
if 'y_pred_xgb' in locals():
    y_pred = y_pred_xgb
    model_label = "XGBoost"
else:
    y_pred = y_pred_rf
    model_label = "Random Forest"

# ---- Get test set metadata ----
# If you already have test_idx from your train_test_split, otherwise get indices another way
try:
    test_idx = X_test.index
except:
    test_idx = customer_df.index[-len(y_test):] # fallback for default split

test_customers = customer_df.loc[test_idx].copy()
test_customers["Actual"] = y_test
test_customers["Predicted"] = y_pred
test_customers["Error"] = test_customers["Predicted"] - test_customers["Actual"]

# 3. Print top 10 largest positive and negative errors (over/under)
largest_pos = test_customers.nlargest(10, "Error")
largest_neg = test_customers.nsmallest(10, "Error")

print(f"{model_label} Model: Actual vs. Predicted Sales Plot")
print(f" Test set size: {len(y_test)} customers")
print(f" Actual sales range: ${y_test.min():.0f} to ${y_test.max():.0f}")
print(f" Predicted sales range: ${y_pred.min():.0f} to ${y_pred.max():.0f}")
print(f" Mean Absolute Error: ${mean_absolute_error(y_test, y_pred):.2f}")
print(f" R2 Score: {r2_score(y_test, y_pred):.3f}\n")

print("Top 10 Largest Positive Prediction Errors (Over-predicted):")
for _, row in largest_pos.iterrows():
    print(f"{row['CustName']:<24} | Predicted: ${row['Predicted']:.0f} | Actual: ${row['Actual']:.0f} | Error: +${row['Error']:.0f}")

print("\nTop 10 Largest Negative Prediction Errors (Under-predicted):")
for _, row in largest_neg.iterrows():
    print(f"{row['CustName']:<24} | Predicted: ${row['Predicted']:.0f} | Actual: ${row['Actual']:.0f} | Error: ${row['Error']:.0f}")

# ---- Plot ----
plt.figure(figsize=(14, 8), facecolor='black')
plt.scatter(
    y_test, y_pred,
    alpha=0.85,
    s=120,
    c='deepskyblue',
    edgecolor='white',
    linewidth=1.4
)

# Reference Line (perfect prediction)
plt.plot(
    [y_test.min(), y_test.max()],
    [y_test.min(), y_test.max()],
    color='lime', linestyle='--', linewidth=3, label="Perfect Prediction"
)

plt.title(f"{model_label}: Actual vs. Predicted Sales", fontsize=28, color='yellow', fontweight='bold', pad=10)
```

```

plt.xlabel("Actual Avg Total Sales per Customer", fontsize=19, color='yellow', fontweight='bold')
plt.ylabel("Predicted Avg Total Sales per Customer", fontsize=19, color='yellow', fontweight='bold')
plt.xticks(fontsize=16, color='white')
plt.yticks(fontsize=16, color='white')
plt.grid(True, linestyle=':', color='gray', alpha=0.4)
plt.gca().set_facecolor('black')
for spine in plt.gca().spines.values():
    spine.set_edgecolor('magenta')
    spine.set_linewidth(2)
plt.legend(facecolor='black', edgecolor='lime', fontsize=16, loc='upper left', labelcolor='white')

# 1: Red arrow for largest positive error
row_pos = largest_pos.iloc[0]
plt.annotate(
    f"Largest Over-Predicted:\n{row_pos['CustName']}",
    xy=(row_pos["Actual"], row_pos["Predicted"]),
    xytext=(row_pos["Actual"] - 225, row_pos["Predicted"] + 500),
    arrowprops=dict(facecolor='red', edgecolor='red', arrowstyle='->', lw=2.7),
    fontsize=14, color='red', fontweight='bold',
    bbox=dict(boxstyle='round', fc='black', ec='red', alpha=0.85)
)

# 2: Red arrow for largest negative error
row_neg = largest_neg.iloc[0]
plt.annotate(
    f"Largest Under-Predicted:\n{row_neg['CustName']}",
    xy=(row_neg["Actual"], row_neg["Predicted"]),
    xytext=(row_neg["Actual"] - 225, row_neg["Predicted"] - 500),
    arrowprops=dict(facecolor='black', edgecolor='red', arrowstyle='->', lw=2.7),
    fontsize=14, color='lime', fontweight='bold',
    bbox=dict(boxstyle='round', fc='black', ec='lime', alpha=0.85)
)

plt.tight_layout()
plt.show()

```

XGBoost Model: Actual vs. Predicted Sales Plot

Test set size: 78 customers

Actual sales range: \$832 to \$3,013

Predicted sales range: \$844 to \$3,262

Mean Absolute Error: \$269.12

R² Score: 0.427


Top 10 Largest Positive Prediction Errors (Over-predicted):

Goldstar Associates	Predicted: \$1,796	Actual: \$1,040	Error: +\$756
Unity Global	Predicted: \$1,842	Actual: \$1,089	Error: +\$753
Sunset Truck Lines	Predicted: \$2,509	Actual: \$1,801	Error: +\$707
Evergreen Corp	Predicted: \$1,523	Actual: \$832	Error: +\$691
Evergreen Global	Predicted: \$1,995	Actual: \$1,381	Error: +\$614
Stonegate Industries	Predicted: \$2,520	Actual: \$1,927	Error: +\$593
Paramount LLC	Predicted: \$2,522	Actual: \$1,950	Error: +\$572
Aurora Partners	Predicted: \$1,528	Actual: \$1,080	Error: +\$448
Stonegate Tucking	Predicted: \$1,564	Actual: \$1,128	Error: +\$436
Summit Solutions	Predicted: \$2,421	Actual: \$2,002	Error: +\$419

Top 10 Largest Negative Prediction Errors (Under-predicted):

Synergy Transfer	Predicted: \$1,487	Actual: \$2,668	Error: \$-1,181
Goldstar Global	Predicted: \$2,104	Actual: \$2,968	Error: \$-864
Sterling Transfer	Predicted: \$1,765	Actual: \$2,266	Error: \$-501
Keystone Logistics	Predicted: \$1,452	Actual: \$1,918	Error: \$-466
Synergy Industries	Predicted: \$1,850	Actual: \$2,303	Error: \$-453
Sunset Tucking	Predicted: \$1,674	Actual: \$2,069	Error: \$-395
Aurora Group	Predicted: \$1,717	Actual: \$2,099	Error: \$-383
Sterling Group	Predicted: \$1,422	Actual: \$1,778	Error: \$-355
Aurora Transfer	Predicted: \$1,402	Actual: \$1,756	Error: \$-354
Sterling Associates	Predicted: \$1,737	Actual: \$2,085	Error: \$-349

XGBoost: Actual vs. Predicted Sales



```
In [34]: from sklearn.metrics import mean_absolute_error, r2_score

# --- Use predictions from your best model (assumes y_pred_xgb, y_test, customer_df, X_test defined) ---
if 'y_pred_xgb' in locals():
    y_pred = y_pred_xgb
    model_label = "XGBoost"
else:
    y_pred = y_pred_rf
    model_label = "Random Forest"

# --- Prepare residuals and test customer metadata ---
residuals = y_test - y_pred

# If you have test_idx from train_test_split, use it; else fallback:
try:
    test_idx = X_test.index
except:
    test_idx = customer_df.index[-len(y_test):]

test_customers = customer_df.loc[test_idx].copy()
test_customers["Actual"] = y_test
test_customers["Predicted"] = y_pred
test_customers["Residual"] = residuals

# Top 10 over- and under-predicted
largest_pos = test_customers.nlargest(10, "Residual")
largest_neg = test_customers.nsmallest(10, "Residual")

print(f"{model_label} Model: Residuals Plot")
print(f"Test set size: {len(y_test)} customers")
print(f"Mean Absolute Error: ${mean_absolute_error(y_test, y_pred):.2f}")
print(f"R2 Score: {r2_score(y_test, y_pred):.3f}\n")

print("Top 10 Largest Positive Residuals (Actual > Predicted):")
for _, row in largest_pos.iterrows():
    print(f"{row['CustName']:<24} | Predicted: ${row['Predicted']:.0f} | Actual: ${row['Actual']:.0f} | Residual: ${row['Residual']:.0f}")

print("\nTop 10 Largest Negative Residuals (Actual < Predicted):")
for _, row in largest_neg.iterrows():
    print(f"{row['CustName']:<24} | Predicted: ${row['Predicted']:.0f} | Actual: ${row['Actual']:.0f} | Residual: ${row['Residual']:.0f}")

# --- Residuals Plot (visual) ---
plt.figure(figsize=(14, 8), facecolor='black')
```

```

plt.scatter(
    y_pred, residuals,
    alpha=0.85, s=120, c='blue',
    edgecolor='white', linewidth=1.0
)
plt.axhline(0, color='lime', linestyle='--', linewidth=3, label="Zero Residual")

# Highlight Largest positive/negative residuals
idx_max = np.argmax(residuals)
idx_min = np.argmin(residuals)
plt.scatter(
    y_pred[idx_max], residuals[idx_max],
    s=240, c='red', edgecolor='yellow', linewidth=2.2, marker='o', zorder=10, label='Largest Over'
)
plt.scatter(
    y_pred[idx_min], residuals[idx_min],
    s=240, c='lime', edgecolor='white', linewidth=2.2, marker='o', zorder=10, label='Largest Under'
)

# Annotate
plt.annotate(
    f"Largest Over: {test_customers.iloc[idx_max]['CustName']}",
    xy=(y_pred[idx_max], residuals[idx_max]),
    xytext=(y_pred[idx_max] - 150, residuals[idx_max] + 400),
    arrowprops=dict(facecolor='red', edgecolor='red', arrowstyle='->', lw=2.2),
    fontsize=14, color='red', fontweight='bold',
    bbox=dict(boxstyle='round', fc='black', ec='red', alpha=0.82)
)
plt.annotate(
    f"Largest Under: {test_customers.iloc[idx_min]['CustName']}",
    xy=(y_pred[idx_min], residuals[idx_min]),
    xytext=(y_pred[idx_min] + 200, residuals[idx_min] + 200),
    arrowprops=dict(facecolor='red', edgecolor='red', arrowstyle='->', lw=2.2),
    fontsize=14, color='lime', fontweight='bold',
    bbox=dict(boxstyle='round', fc='black', ec='lime', alpha=0.82)
)

plt.title(f"{model_label}: Residuals Plot", fontsize=28, color='darkorange', fontweight='bold', pad=10)
plt.xlabel("Predicted Avg Total Sales per Customer", fontsize=22, color='darkorange', fontweight='bold')
plt.ylabel("Residual (Actual - Predicted)", fontsize=22, color='darkorange', fontweight='bold')
plt.xticks(fontsize=16, color='white')
plt.yticks(fontsize=16, color='white')
plt.grid(True, linestyle=':', color='gray', alpha=0.3)
plt.gca().set_facecolor('black')
for spine in plt.gca().spines.values():
    spine.set_edgecolor('orange')
    spine.set_linewidth(2)
plt.legend(facecolor='black', edgecolor='yellow', fontsize=19, loc='upper right', labelcolor='white')
plt.tight_layout()
plt.show()

```

XGBoost Model: Residuals Plot

Test set size: 78 customers

Mean Absolute Error: \$269.12


R² Score: 0.427

Top 10 Largest Positive Residuals (Actual > Predicted):

Synergy Transfer	Predicted: \$1,487	Actual: \$2,668	Residual: +\$1,181
Goldstar Global	Predicted: \$2,104	Actual: \$2,968	Residual: +\$864
Sterling Transfer	Predicted: \$1,765	Actual: \$2,266	Residual: +\$501
Keystone Logistics	Predicted: \$1,452	Actual: \$1,918	Residual: +\$466
Synergy Industries	Predicted: \$1,850	Actual: \$2,303	Residual: +\$453
Sunset Tucking	Predicted: \$1,674	Actual: \$2,069	Residual: +\$395
Aurora Group	Predicted: \$1,717	Actual: \$2,099	Residual: +\$383
Sterling Group	Predicted: \$1,422	Actual: \$1,778	Residual: +\$355
Aurora Transfer	Predicted: \$1,402	Actual: \$1,756	Residual: +\$354
Sterling Associates	Predicted: \$1,737	Actual: \$2,085	Residual: +\$349

Top 10 Largest Negative Residuals (Actual < Predicted):

Goldstar Associates	Predicted: \$1,796	Actual: \$1,040	Residual: \$-756
Unity Global	Predicted: \$1,842	Actual: \$1,089	Residual: \$-753
Sunset Truck Lines	Predicted: \$2,509	Actual: \$1,801	Residual: \$-707
Evergreen Corp	Predicted: \$1,523	Actual: \$832	Residual: \$-691
Evergreen Global	Predicted: \$1,995	Actual: \$1,381	Residual: \$-614
Stonegate Industries	Predicted: \$2,520	Actual: \$1,927	Residual: \$-593
Paramount LLC	Predicted: \$2,522	Actual: \$1,950	Residual: \$-572
Aurora Partners	Predicted: \$1,528	Actual: \$1,080	Residual: \$-448
Stonegate Tucking	Predicted: \$1,564	Actual: \$1,128	Residual: \$-436
Summit Solutions	Predicted: \$2,421	Actual: \$2,002	Residual: \$-419



Feature Importance

We visualize which features most influence the model's predictions.

```
In [35]: # Determine which model to use for feature importance
```

```
if 'xgb' in locals():
    importances = xgb.feature_importances_
    feature_names = X.columns
    model_name = "XGBoost"
elif 'rf' in locals():
    importances = rf.feature_importances_
    feature_names = X.columns
    model_name = "Random Forest"
else:
```

```

importances = None
feature_names = []

if importances is not None:
    # Get top N features for readability
    N = 15
    indices = np.argsort(importances)[::-1][:N]
    top_features = [(feature_names[i], importances[i]) for i in indices]

    # Print the top 10 features with importances
    print(f"{model_name} Model: Top 10 Feature Importances")
    for i, (feat, imp) in enumerate(top_features[:10], 1):
        print(f"{i:2d}. {feat:<35} | Importance: {imp:.4f}")
    print()

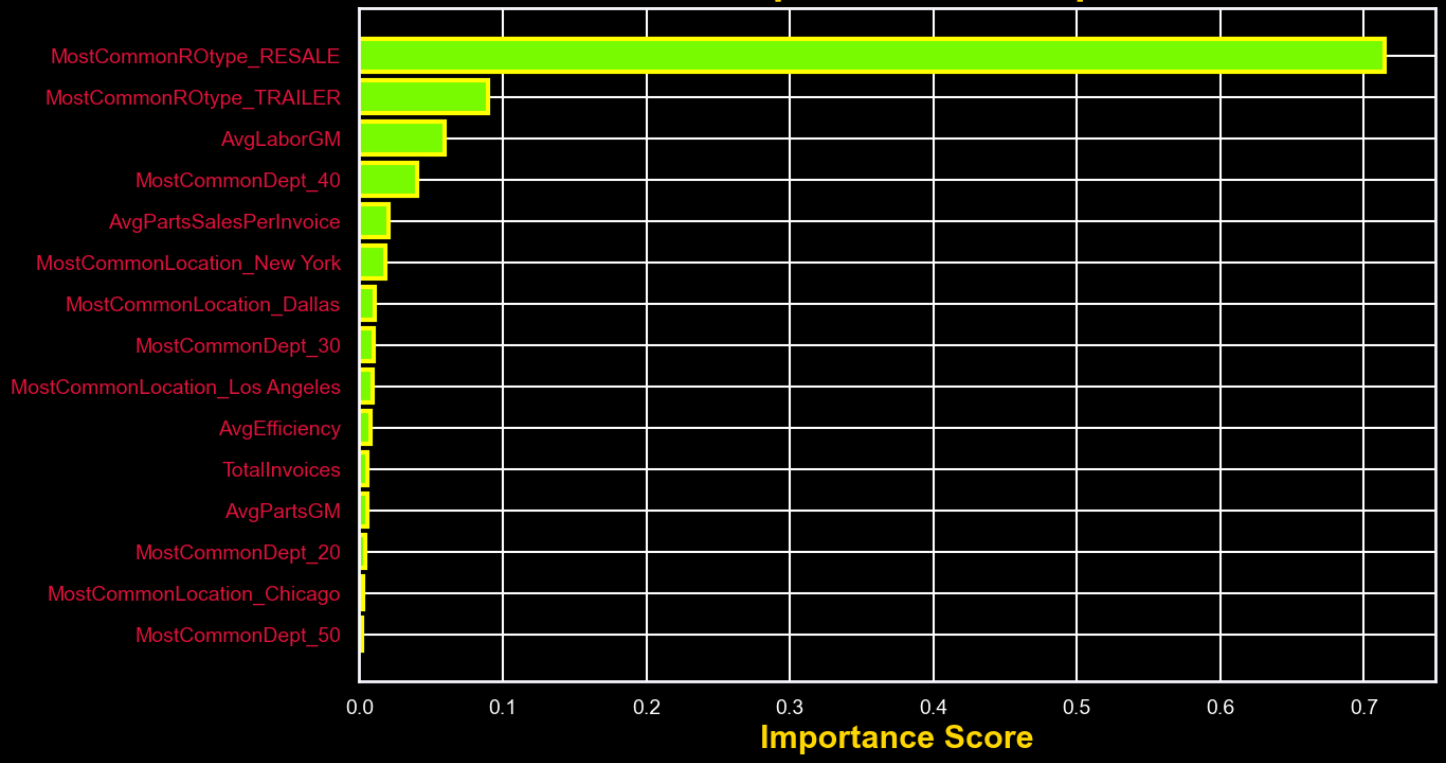
    # Plot
    plt.figure(figsize=(14, 8), facecolor='black')
    bars = plt.barh(
        range(N),
        [importances[i] for i in indices],
        color='lawngreen',
        edgecolor='yellow',
        linewidth=3
    )
    plt.yticks(range(N), [feature_names[i] for i in indices], fontsize=14, color='crimson')
    plt.gca().invert_yaxis()
    plt.xlabel('Importance Score', fontsize=22, color='gold', fontweight='bold')
    plt.title(f'{model_name}: Top {N} Feature Importances', fontsize=28, color='gold', fontweight='bold', pad=10)
    plt.xticks(fontsize=14, color='white')
    plt.gca().set_facecolor('black')
    for spine in plt.gca().spines.values():
        spine.set_edgecolor('ghostwhite')
        spine.set_linewidth(2)
    plt.tight_layout()
    plt.show()
else:
    print("Feature importance not available.")

```

XGBoost Model: Top 10 Feature Importances

1. MostCommonR0type_RESALE	Importance: 0.7141
2. MostCommonR0type_TRAILER	Importance: 0.0900
3. AvgLaborGM	Importance: 0.0593
4. MostCommonDept_40	Importance: 0.0404
5. AvgPartsSalesPerInvoice	Importance: 0.0201
6. MostCommonLocation_New York	Importance: 0.0178
7. MostCommonLocation_Dallas	Importance: 0.0105
8. MostCommonDept_30	Importance: 0.0099
9. MostCommonLocation_Los Angeles	Importance: 0.0090
10. AvgEfficiency	Importance: 0.0080

XGBoost: Top 15 Feature Importances



Conclusion

This project built and evaluated multiple regression models to predict average total sales per dealership customer using operational and categorical data. The analysis confirmed that customer segmentation, especially by job type (RESALE vs COUNTER), is the strongest driver of sales. Simpler models like linear regression performed just as well as advanced methods for this dataset, with an R^2 around 0.57 and a mean absolute error near \$236.

The XGBoost model, while robust, did not improve accuracy, likely due to the linear nature of business relationships in this data. Feature importance analysis reinforced that customer type, labor gross margin, and specific departments or locations are most predictive of customer value.

Business teams should use these insights to focus on growing the RESALE segment, increasing service work among parts-heavy customers, and monitoring key accounts for changes in purchasing behavior. Outlier accounts identified by the model offer immediate opportunities for review, engagement, or risk management.

With more real-world data or additional features, model accuracy may improve further. This project provides a solid foundation for data-driven customer segmentation and growth planning in dealership operations.